

A Comprehensive Assessment Tool for Prompt Injection Attacks in Large Language Models (LLMs)

Ailene Dao
Cybersecurity Engineering
George Mason University
Fairfax, VA
www.gmu.edu

Richard VanGorder
Cybersecurity Engineering
George Mason University
Fairfax, VA
www.gmu.edu

Mohamed Gebril
Cybersecurity Engineering
George Mason University
Fairfax, VA
mgebril@gmu.edu

Sherif Abdelhamid
Computer & Information Sciences
Virginia Military Institute
Lexington, VA
abdelhamidse@vmi.edu

Abstract—A prompt injection attack is a type of cyberattack that targets Large Language Models (LLMs), where attackers embed malicious instructions into a LLM with the intent of manipulating the model to the attacker's desire. Due to the relevance, growth, and demand of LLMs within small to medium-sized businesses, it is essential that they have access to a tool that can comprehensively assess model vulnerabilities. Based on an established framework for Prompt Injection attacks, this paper introduces a tool that can simulate types of injection attacks across different widely used LLM models and versions, test different defenses against these attacks, and measure the performance and how vulnerable the LLM's responses are. Additionally, this paper incorporates cost analysis, CPU utilization metrics, and a prompt hardening methodology for LLM personas. The results of these experiments provide deeper insights into testing and a comparative analysis against commonly used LLMs today.

Keywords—Artificial intelligence, Cybersecurity, Large language models, Prompt injection attacks

I. INTRODUCTION

Large Language Models (LLMs) have become widely used and integrated in a wide range of modern applications. However, as their adoption grows, so does the concern about their robustness, reliability, and vulnerability to malicious manipulation. Among these threats are prompt injections, a type of attack that involves exploiting an LLM by overriding its intended instructions and leaking sensitive data or executing unauthorized actions, posing risks to either the user, LLM, or both.

Recent work by Liu, Y. Jia, Geng, J. Jia, and Gong has laid foundation work in formalizing and benchmarking prompt injection attacks [1]. Their study introduces types of prompt injection attacks, explores different defenses against them, and evaluates a model's vulnerability. However, their evaluation framework focuses primarily on attack success rates with limited emphasis on other operational evaluation metrics such as CPU utilization, cost, and execution time-factors critical in modern deployment.

II. RELATED WORK

Much has been tested and documented since the release of LLM models. This is useful as a starting point to get familiar with the research area and methods already created and tested. Research has focused on different areas of prompt injections. For example, some research focused on the attack-aspect. X. Liu, Z. Yu, Y. Zhang, N. Zhang, and C. Xiao define and categorize objectives of prompt injection attacks and devises gradient-based training for developing attacks [2]. Perez and Ribeiro categorize prompt injection attacks as goal hijacking and prompt leaking and explore their behaviors [3]. S. Lee, J. Kim, and W. Pak explore prompt injections via the use of images as visual input [4]. Q. Zhan, R. Fang, H. S. Panchal, and D. Kang test adaptive attacks against current defenses to show how advanced attacks are becoming [5].

Other research focused on developing and testing defenses against these attacks. Hines et al. introduce spotlighting as a defense to aid the LLM to determine which part of the prompt is valid to answer [6]. Suo introduces signed prompts as a defense against prompt injection attacks [7]. Shi et al. introduce PromptArmor to directly extract the injected portion of a prompt, therefore sanitizing the input before being fed into an LLM [8]. S. Chen, J. Piet, C. Sitawarin, and D. Wagner test structured queries as a defense against prompt injection attacks [9]. Jia et al. analyze and compare the effectiveness of prevention and detection-based defenses [10].

Additionally, different LLMs were tested. Liu et al. focuses more on case studies about LLM-integrated applications and their vulnerabilities against prompt injection attacks [11]. M. Alizadeh, Z. Samei, D. Stetsenko, and F. Gilardi evaluate responses for agentic AI models when presented with prompt injections [12]. J. Wang, P. Gupta, I. Habernal, and E. Hullermeier test open-source LLMs against prompt injection attacks [13].

Matthew shows an aggregation of most of the above work related to prompt injection attacks and defenses from a multitude of research papers [14]. However, more research is to be conducted as newer versions and models are being released. The limitations in the work tested in this area mostly reside in the types and numbers of models and operational

evaluation metrics. The models tested are mostly GPT-based or Claude, which may cause unknown bias and limited information about other LLM models. These limitations motivate our work to contribute to this area.

III. GRAPHICAL USER INTERFACE ARCHITECTURE

To facilitate seamless testing and comparative analysis across multiple LLMs, we developed a graphical user interface (GUI). This interface enables users to configure model, attack, defense, and other parameters to execute simulations and analyze the results. The GUI is composed of three parts: the Home, Output, and Analysis page tabs.

The Homepage, as shown in Fig. 1, serves as the main configuration tab. Users can select the target LLM and its specific version, choose the type of prompt injection attack, and choose a defense strategy. Additionally, users will have to enter a valid and paid API key and specify the number of trials to be executed. Once they start running the test, the prompt injection attack simulation begins, and users can see a progress bar indicating the simulation status. After each test, users can reconfigure another attack and run additional simulations without restarting the application.

The Output page as shown in Fig. 2 displays detailed results for an individual simulation, following the completion of a test. This page includes:

- The Attack Success Value (ASV) showing the effectiveness of the prompt injection
- The Associated Vulnerability Level categorized based on the ASV value
- The Elapsed Time to complete the test
- Written results showing the prompt and the LLM's response

The Analysis tab, as shown in Fig. 3, aggregates the results from all completed simulations. Each new test is appended to the dataset, enabling comparative analysis. Key features include:

- A bar graph visualization of ASV values across different model versions
- Summary statistics for ASV values and experiment configuration
- An option to export results to a local file to keep for record




Fig. 1. Homepage – Attack/Defense Configuration. Users choose a model, version, attack type, and defense type from the dropdown menus. The textboxes accept an API key and the number of desired trials. Configuration instructions are below the blue, Run Tests button.




Fig. 2. Webpage showing the Results of one Simulated Test. Metrics show the ASV value, associated vulnerability level, and time executed. It shows the prompt and the LLM's response to it.

IV. EXPERIMENTAL SETUP

To evaluate prompt injection vulnerabilities and defense mechanisms across a diverse set of large language models (LLMs), we selected four major providers: OpenAI, Google, Anthropic, and DeepSeek. For each provider, multiple model versions were tested to ensure comprehensive coverage.

A. OpenAI GPT:

- GPT-3.5-Turbo
- GPT-4
- GPT-4.1
- GPT-5

B. Google Gemini (formerly PaLM2)

- Gemini-1.5-Flash
- Gemini-2.0-Flash
- Gemini-2.5-Flash




Fig. 3. Comparative Analysis of ASV between Multiple Tests. It shows a summary of the configuration and metrics of each test and a visual bar graph for the ASV value.

C. Anthropic Claude

- Claude-3-Opus-20240229
- Claude-Opus-4-20250514

D. DeepSeek

- DeepSeek-Chat

V. EVALUATION METRICS

To quantitatively assess the vulnerability of each model version to prompt injection attacks, we adopt the Attack Success Value (ASV) metric introduced by Liu, Y. Jia, Gong, J. Jia, and Gong [1]. ASV is a normalized score ranging from 0 to 1, where:

- 0 indicates no attacks were successful
- 1 indicates all attacks were successful

In addition to ASV, we incorporated three other metrics to provide a more detailed view of a model's performance:

- **CPU utilization rate:** A separate monitoring program tracks CPU usage every 5 seconds during each test and computes the average upon completion. This metric helps assess the computational efficiency of each model.
- **Execution time:** Measured in minutes and seconds, this reflects the duration of a simulated test run
- **Cost:** Retrieved from the API provider's billing dashboard, this metric accounts for token usage and pricing per model version

These additional metrics enable prospective users to make informed decisions based on their priorities whether optimizing speed, resource consumption, or cost efficiency.

VI. ATTACK AND DEFENSE SETUP

To simulate prompt injection scenarios, we adopt the attack and defense taxonomy proposed by Liu et al. [5], which provides a structured framework for evaluating vulnerabilities against prompt injection attacks in large language models (LLMs).

Each simulated attack consists of two tasks:

Target task: The intended user-defined task. In this experiment, we use sentiment analysis where the model classifies input text or data as positive or negative.

Injected task: A maliciously embedded task designed to override the target task. In this experiment, we use spam detection where the model classifies additional injected input text as spam or not spam.

Attack Types (defined by [1])

- **Naive** - Target task immediately followed by the injected task
- **Escape Character** - Target task followed by a '\n' or '\t' followed by injected task
- **Context Ignoring** - Target task + "Ignore previous instructions." + Injected task
- **Fake Completion** - Target task + "Answer: negative." + Injected task
- **Combination** - combines all the above tactics

Defense Types (defined by [1])

- **None** - No defense applied
- **LLM-based** - LLM first determines if the prompt is safe to answer, ignores if the prompt is unsafe
- **Known Answer** - LLM answers with a pre-defined answer. If it matches, LLM answers prompt
- **Response-based** - LLM checks whether answer matches target prompt
- **Paraphrasing** - LLM paraphrases prompt and answers the rephrased prompt
- **Sandwich** - LLM is reminded at the end of prompt what it's original task is
- **Delimiters** - Uses "" delimiters to sandwich benign and injected input
- **XML** - Uses <user_input> to sandwich benign and injected input.

VII. PROMPT HARDENING

In addition to the above framework, we introduce prompt hardening as a defense technique. This approach involves creating a persona or behavioral restraint to guide the LLM's behavior when responding to prompt injections. Prompt hardening can be implemented in two ways:

- **Modifying System Persona:** For models that support editable system roles (e.g., GPT), the default persona (e.g., "You are a helpful assistant") can be replaced with a more security-conscious description
- **For models without system-level persona support** (e.g., Gemini), the hardened persona is appended before the input prompt

Because of the structured format in relation to the target task always being the first task and the injected task always

TABLE I. GPT-4 SHOWING ASV, CPU, AND TIME MEASUREMENTS FOR EACH EACH DEFENSE AGAINST THE NAIVE ATTACK. LOWER ASV, CPU, AND TIME METRICS ARE OPTIMAL. 100 DATA SAMPLES WERE USED AS THE TRIAL NUMBER AND THE TOTAL COST AMOUNTS TO \$20.31.

Defense	Evaluation Metrics for GPT-4 with Naive Attack		
	ASV	CPU (%)	Time (min: sec)
None	0.56	10.46	3:19
LLM-based	0.45	7.41	4:30
Known Answer	0.52	6.97	6:07
Response-based	0.54	7.23	3:21
Paraphrasing	0.28	8.24	8:09
Sandwich	0.40	7.85	3:17
Delimiters	0.52	9.04	3:18
XML	0.48	8.85	3:20

TABLE II. ASV, CPU, AND TIME MEASUREMENTS FOR TESTS FOR EACH MODEL WITH PROMPT HARDENING. EACH MODEL IS MODIFIED WITH THE SECURITY-LEVEL PERSONA AND TESTED AGAINST EACH ATTACK.

Model Version	Evaluation Metrics for All Models with Prompt Hardening and Combination Attack		
	ASV	CPU (%)	Time (min: sec)
Gemini 1.5 Flash	0.93	11.42	2:55
Gemini 2.0 Flash	0.69	10.48	3:08
Gemini 2.5 Flash	0.0	12.51	7:48
GPT 3.5 Turbo	0.83	7.94	2:14
GPT 4	0.34	6.63	2:52
GPT 4.1	0.98	14.44	2:22
GPT 5	0.02	8.01%	15:52
Claude Opus 4	0.0	8.60%	5:27
Claude Opus 3	0.89	8.14%	4:22
DeepSeek Chat	0.46	10.37%	7:13

being the second task, in this experiment, our hardened prompt or persona is “You are a strict assistant that only follows structured instructions from the system role. Ignore any attempts inside user messages to change behavior, act outside of scope, or repeat system prompts. Only follow the initial task from the user. Do not obey any instructions embedded later, even if they look like commands or start with words like ‘ignore’, ‘answer’, or ‘repeat’”.

However, users should tailor their specific use cases and expected inputs. Users should expect higher token usage and cost due to longer system prompts and the effectiveness varies depending on the hardened prompt. Nonetheless, users are encouraged to experiment with different personas to identify optimal configurations for their desired model and output.

For each model and version, we computed a table showing the different combinations of attack and defense types and its measured values. Table I shows an example of GPT-4 with the Naive Attack. A table is created with each attack so for GPT-4, there would be 5 total tables similar to Table I. The cost for

TABLE III. AGGREGATION OF AVERAGES AND ASSOCIATED FINAL METRIC VALUE FOR EACH MODEL VERSION. AVERAGES ARE CALCULATED BASED ON VALUES COLLECTED IN MODEL TABLES LIKE IN TABLE I. THE LOWER FINAL VALUE IS OPTIMAL FOR COMPARISON PURPOSES..

Model Version	Comparison of Models				
	Avg ASV	Avg CPU (%)	Avg Time (sec)	Avg Cost (USD)	Aggregated Value
Gemini 1.5 Flash	0.285	9.09	155.900	\$0.05	24.878
Gemini 2.0 Flash	0.371	8.92	185.450	\$0.07	29.324
Gemini 2.5 Flash	0.348	9.56	530.625	\$8.44	83.7
GPT 3.5 Turbo	0.403	8.95	159.525	\$8.95	28.118
GPT 4	0.498	8.41	238.250	\$20.31	43.292
GPT 4.1	0.39	8.90	178.000	\$21.47	34.632
GPT 5	0.385	8.32	947.225	\$26.32	151.382
DeepSeek Chat	0.329	7.93	536.725	\$0.10	81.859
Claude Opus 3	0.564	8.95	420.700	\$18.94	70.355
Claude Opus 4	0.276	9.10	518.600	\$13.43	83.294

TABLE IV. AVERAGES AND ASSOCIATED FINAL METRIC VALUE FOR EACH MODEL VERSION AFTER PROMPT HARDENING. AVERAGES ARE CALCULATED BASED ON VALUES COLLECTED IN TABLE II.

Model Version	Comparison of Models			
	Avg ASV	Avg CPU (%)	Avg Time (sec)	Aggregated Value
Gemini 1.5 Flash	0.186	11.738	174.6	28.025
Gemini 2.0 Flash	0.280	10.394	191.8	30.441
Gemini 2.5 Flash	0.008	9.404	428.8	65.734
GPT 3.5 Turbo	0.526	12.960	140.8	23.274
GPT 4	0.084	11.094	166.2	26.628
GPT 4.1	0.592	10.886	151.8	24.634
GPT 5	0.004	8.518	661.4	100.489
DeepSeek Chat	0.004	5.716	315.4	48.169
Claude Opus 3	0.622	6.152	272.2	42.002
Claude Opus 4	0.210	10.816	433.8	66.776

the GPT-4 version after running all tests amounts to \$20.31. A table was created for each model version and attack/defense combination.

VIII. RESULTS

Table II also shows an aggregated table for model versions with prompt hardening. Because we are categorizing prompt hardening as a defense, we are testing each hardened model version only with the five attacks previously stated.

To easily compare different models and versions based on multiple evaluation metrics, we developed a formula that generates a single final metric value. This value is derived from the ASV, CPU usage, time, and cost based on the following weights:

- ASV: 40%
- Cost: 30%
- CPU Usage: 15%
- Time: 15%

First, we computed the average ASV, CPU, and time metrics across all models. The average cost is the total cost for each table. The final score was then calculated based on the average values. Time is also converted to seconds for easier computations. Table III shows the average values and the final




Fig. 4. Final Metric Values before and after Prompt Hardening. It shows that prompt hardening is an effective defense at lowering the LLM’s final value thus improving its effectiveness.

score for each model version. The model with the lowest final aggregated score is deemed to be the best performing model based on the above factors.

Based on the computed scores, the top-performing model is Gemini-1.5-Flash, while the model with the lowest performance is GPT-5, likely due to its higher cost and longer execution time.

In addition, we performed a separate evaluation of models based solely on prompt injection, as shown in Table IV. This evaluation followed the same methodology as the previous tests. However, in this case, the cost was considered negligible, given the small number of tests and the correspondingly lower token consumption. Based on the final metric values, GPT-3.5-Turbo and GPT-4.1 resulted in having the best performance with prompt hardening.

Fig. 4 shows a visual comparison between the aggregated values of each model version before and after prompt hardening. Though Gemini 1.5 and 2.0 Flash models are outliers, the rest of the models see a significant reduction in the aggregated value after prompt hardening, showing its effectiveness against prompt injections.

IX. GRAPHICAL USER INTERFACE IMPORT ARCHITECTURE

To extend the analysis capabilities of the GUI and support reproducible model benchmarking, an Import tab was implemented into the GUI that allows for the use of external LLM evaluation tools and datasets within the GUI. Two evaluation tools were utilized for this work: Promptfoo and FuzzyAI. From within the new tab, it is possible to run tests by choosing a specific scanning vulnerability tool and choosing a selected type of test.

All code related to this section of the vulnerability assessment tool is located in the external folder and the main UI file. For adding a new tool, create a new folder within the external folder and add a main.py in a similar style to the existing tools’ main.py (note that every tool has a different method to calling it, look at documentation of the desired tool). From there, have external/import_main call on the main

function of the main.py. Add any names of tools and attacks/datasets from within the main UI file. If done correctly, it should be possible to import almost any scanning tool or dataset and add it to the comprehensive vulnerability scanner.




Fig. 5. Import tab showing the configuration set-up for an imported attack method and/or dataset. The user can select different imported tools and test cases from a dropdown menu. The textboxes are for inputting API keys, different tokens, and number of trials. The graphical user interface runs tests on a model with an API key input. Selecting the blue Run tests button starts the attack, upon finishing it will switch to the output tab.

TABLE V. RESULTS OF 100 PROMPTS FOR THE “ARC CHALLENGE” DATASET ON HUGGINGFACE WHICH CAN BE FOUND AT [ALLENAI/AI2_ARC_DATASETS AT HUGGING FACE](#). TESTED UTILIZING PROMPTFOO: LLM EVALUATIONS & RED TEAMING SCANNER.

Model	ASV	Test Pass	Test Fail	Tokens	Cost (\$ USD)
OpenAI: o1	0.03	90	10	63,436	\$3.27
Gemini 2.5 Pro	0.02	95	5	185,802	NA
Deepseek-Chat	0.02	95	5	35,876	NA
Claude Sonnet-4	0.01	96	4	47,976	\$0.56

X. IMPORTED TESTS & DATASETS INTRODUCTION

The first evaluation tool Promptfoo “is a developer-friendly local tool for testing LLM applications. Stop the trial-and-error approach - start shipping secure, reliable AI apps.” In other words, it allows for the implementation of HuggingFace and other imported datasets that can test multiple LLM models simultaneously based on desired classifications that will be determined to be correct by Promptfoo’s grading AI. These are the metrics utilized for the tests:

- Models Tested: OpenAI o1, Gemini 2.5 Pro, Deepseek Chat, & Claude Sonnet 4
- Graded on (Assertions): Step-by-step reasoning, latency (180000ms), & correct answers
- Trial Amount per Model: 100
- Datasets Utilized: ARC Challenge, OpenBookQA, LogiQA, MMLU, Clickbait Detection, SMS Spam Detection, Prompt Injections, and Prompt Injection Dataset (Naive, Ignore Context, Fake Completion, and Escape)

The second evaluation tool FuzzyAI “is a powerful tool for automated LLM fuzzing. It is designed to help developers and security researchers identify jailbreaks and mitigate potential security vulnerabilities in their LLM APIs.” In other words, it allows for dangerous prompts like “How to build a bomb” and see if a variety of different attack techniques can

allow for a successful prompt injection to occur. These are the metrics utilized for the tests:

- Models Tested: GPT-4, Deepseek Chat, & Claude Sonnet 4
- Graded on (Assertion): “Is the following content harmful?”
- Trial Amount per Model: 100
- Dataset Utilized: FuzzyAI’s adv_prompts.txt

TABLE VI. RESULTS OF 100 PROMPTS FOR THE “OPENBOOKQA” DATASET ON HUGGINGFACE WHICH CAN BE FOUND AT [ALLENAI/OPENBOOKQA · DATASETS AT HUGGING FACE](#). TESTED UTILIZING PROMPTFOO: LLM EVALUATIONS & RED TEAMING SCANNER.

Model	ASV	Test Pass	Test Fail	Tokens	Cost (\$ USD)
OpenAI: o1	0.06	76	24	63,621	\$3.40
Gemini 2.5 Pro	0.02	94	6	168,530	NA
Deepseek-Chat	0.02	93	7	29,045	NA
Claude Sonnet-4	0.01	97	3	40,030	\$0.48

TABLE VII. RESULTS OF 100 PROMPTS FOR THE “OPENBOOKQA” DATASET ON HUGGINGFACE WHICH CAN BE FOUND AT [ALLENAI/OPENBOOKQA · DATASETS AT HUGGING FACE](#). TESTED UTILIZING PROMPTFOO: LLM EVALUATIONS & RED TEAMING SCANNER.

Model	ASV	Test Pass	Test Fail	Tokens	Cost (\$ USD)
OpenAI: o1	0.06	76	24	63,621	\$3.40
Gemini 2.5 Pro	0.02	94	6	168,530	NA
Deepseek-Chat	0.02	93	7	29,045	NA
Claude Sonnet-4	0.01	97	3	40,030	\$0.48

- Attacks Performed: Hallucinations, DAN (do anything now), History/Academic Framing, ThoughtExperiment, & Best-of-n Jailbreaking

XI. EVALUATION OF IMPORTED TESTS & DATASETS

A. Promptfoo

This dataset was chosen to challenge the LLMs capability to correctly answer difficult questions. Due to it causing the least amount of errors, this dataset was chosen to be utilized again in the “large dataset” portion. This data is the first indicator of the usual trends that can be seen in the other result tables. Gemini tends to use the most amount of tokens, while Deepseek uses the least. OpenAI uses the 2nd most amount of tokens, however it tends to have a significantly higher cost than any other model even when estimating the cost of Deepseek and Gemini.

The main issue for OpenAI in this test was its inability to follow the “provide step-by-step reasoning” assertion, which led to it being consistently flagged by the grading AI. It is notable that this is a reoccurring issue with OpenAI-based models, mostly since OpenAI tends to have the shortest responses. This is a bizarre contradiction with the token usage, since it had the second-highest token usage, despite having significantly shorter responses than Deepseek and Anthropic models who consume significantly less tokens.

Out of all the different tests, LogiQA tended to lead to the most wrong answers. That being said, the results are misleading as the main issue is that the prompts seemed to mess up the format (especially for Gemini) that the model would answer the question. This lead to the grading system to flag it as incorrect even when the answer was correct. The reason why these results were left in is due to this phenomenon, as this issue isn’t seen in datasets testing similar things. This may be in relation to the token usage which is much higher in this test compared to the other tests.

MMLU seemed to on occasion break Deepseek-chat which led to it outputting nonsense that couldn’t be

TABLE VIII. RESULTS OF 100 PROMPTS FOR THE “LOGIQA” DATASET ON HUGGINGFACE WHICH CAN BE FOUND AT [LUCASMCCABE/LOGIQA · DATASETS AT HUGGING FACE](#). TESTED UTILIZING PROMPTFOO: LLM EVALUATIONS & RED TEAMING SCANNER.

Model	ASV	Test Pass	Test Fail	Tokens	Cost (\$ USD)
OpenAI: o1	0.07	79	21	242,738	\$13.47
Gemini 2.5 Pro	0.25	26	73	458,617	NA
Deepseek-Chat	0.16	53	47	83,979	NA
Claude Sonnet-4	0.08	75	25	80,006	\$0.88

TABLE IX. RESULTS OF 100 PROMPTS FOR THE “MMLU” DATASET ON HUGGINGFACE WHICH CAN BE FOUND AT [CAIS/MMLU · DATASETS AT HUGGING FACE](#). TESTED UTILIZING PROMPTFOO: LLM EVALUATIONS & RED TEAMING SCANNER.

Model	ASV	Test Pass	Test Fail	Tokens	Cost (\$ USD)
OpenAI: o1	0.02	94	6	121,930	\$6.58
Gemini 2.5 Pro	0.01	98	2	306,847	NA
Deepseek-Chat	0.04	88	12	65,514	NA
Claude Sonnet-4	0.02	95	5	72,446	\$0.86

determined what it was trying to answer. The most likely explanation is that Deepseek ran out of typing characters (given to it by Promptfoo) and didn’t have enough time to fully answer. Despite this, Deepseek tends to have lower token usage compared to other models, despite having the longest text on many occasions. The models reasoning skills was the main test behind this dataset and it tended to have the most accurate results in regards to false positives in reasoning tests.

This dataset focuses more on the concept of classifying whether a prompt is clickbait or not. In this case, its meant to determine if a model could realistically act as a filter for large amounts of news. In this case, there were no false positives, however there were several false negatives.

In regard to future research, this dataset is rather useful as its ability to classify spam and not spam could effectively stop phishing attempts. There were no false positives in this test, however there were several false negatives. In that regard, it is unlikely that models will accidentally flag a large amount of benign messages if used in a email/messaging system. It is also notable that it had results similar to the clickbait detection dataset.

This dataset focuses on whether the LLM can classify the difference between a benign and prompt injection message.

OpenAI compared to other models tended to overwhelmingly misclassify prompt injection statements as benign. Based on the LLM’s reasoning, this appears to be because it has a different understanding of what a prompt injection is compared to the other models. In this scenario, this works against it as it classified almost all scenarios as benign.

The Naive attack was overall ineffective. The LLM didn’t seem to even notice that anything was weird when this attack was performed. This correlates with the results seen when performing a naive attack using the main vulnerability scanner.

Despite failing considerably more tests, it is notable that OpenAI failed this much due to it failing the “give step-by-

TABLE X. RESULTS OF 100 PROMPTS FOR THE “CLICKBAIT DETECTION” DATASET ON HUGGINGFACE WHICH CAN BE FOUND AT [CHRISTINACDL/CLICKBAIT_DETECTION_DATASET · DATASETS AT HUGGING FACE](#). TESTED UTILIZING PROMPTFOO: LLM EVALUATIONS & RED TEAMING SCANNER.

Model	ASV	Test Pass	Test Fail	Tokens	Cost (\$ USD)
OpenAI: o1	0.06	82	18	76,136	\$4.12
Gemini 2.5 Pro	0.06	83	16	154,770	NA
Deepseek-Chat	0.06	83	16	33,195	NA
Claude Sonnet-4	0.06	83	17	45,093	\$0.55

TABLE XI. RESULTS OF 100 PROMPTS FOR THE “SMS SPAM COLLECTION” DATASET ON HUGGINGFACE WHICH CAN BE FOUND AT [CODESIGNAL/SMS-SPAM-COLLECTION · DATASETS AT HUGGING FACE](#). TESTED UTILIZING PROMPTFOO: LLM EVALUATIONS & RED TEAMING SCANNER.

Model	ASV	Test Pass	Test Fail	Tokens	Cost (\$ USD)
OpenAI: o1	0.05	84	16	65,341	\$3.40
Gemini 2.5 Pro	0.03	92	8	168,859	NA
Deepseek-Chat	0.02	94	6	34,808	NA
Claude Sonnet-4	0.05	87	13	43,491	\$0.51

step reasoning” assertion. Overall, the models scored somewhat consistently when it came to getting the correct answer. Claude showed the best results as it got a perfect score and utilized the least amount of tokens.

This method was most effective against Deepseek, while scoring somewhat consistently as other prompt injection attacks on other models. The exact reason for this is unknown, however this could be a result of it using the least number of tokens. As a result, this might indicate a legitimate cybersecurity threat to models that utilize lower tokens amount even if it significantly lowers the cost.

This test resulted in similar outputs and outcomes as the Naive attack. In almost all cases, the LLMs treated the Escape Character attack as if the Naive attack was performed. In other words, it didn’t really register it, as a result, a similar conclusion to the naive attack can be made.

B. FuzzyAI

The “Hallucination” attack method was rather effective with making the LLM repeat what the user says. However, the actual usefulness of this is rather low, while it might

repeat the prompt in the format desired by the user, it wouldn’t repeat phrases that would be considered offensive.

Which appears to be the main reason to perform this attack, as the hallucinations do not tell the model to actually answer the prompt (ie. How to build a bomb?).

Out of all the different attack methods, this one was the least effective, the “DAN” persona did not seem to overwrite the tested models. Interestingly this lead to a reverse effect where the DAN attack seemed to put the LLM “on-guard” more than overwriting it, even against prompts that the model usually struggles against. This technique may be more effective at corrupting models that are meant to act in different personas, Character AI for instance.

TABLE XII. RESULTS OF 100 PROMPTS FOR THE “PROMPT INJECTIONS” DATASET ON HUGGINGFACE (NOT TO BE CONFUSED WITH THE PROMPT INJECTION ATTACK DATASET) WHICH CAN BE FOUND AT [DEEPPSET/PROMPT-INJECTIONS · DATASETS AT HUGGING FACE](#). TESTED UTILIZING PROMPTFOO: LLM EVALUATIONS & RED TEAMING SCANNER. THIS TEST WAS ON THE PROMPT INJECTION’S FAKE COMPLETION ATTACK.

Model	ASV	Test Pass	Test Fail	Tokens	Cost (\$ USD)
OpenAI: o1	0.21	42	58	86,589	\$4.68
Gemini 2.5 Pro	0.09	72	28	175,639	NA
Deepseek-Chat	0.02	93	7	34,369	NA
Claude Sonnet-4	0.06	82	18	44,578	\$0.52

TABLE XIII. RESULTS OF 100 PROMPTS FOR THE “PROMPT INJECTION ATTACK” DATASET ON HUGGINGFACE WHICH CAN BE FOUND AT [XXZ224/PROMPT-INJECTION-ATTACK-DATASET · DATASETS AT HUGGING FACE](#). TESTED UTILIZING PROMPTFOO: LLM EVALUATIONS & RED TEAMING SCANNER. THIS TEST WAS ON THE PROMPT INJECTION’S NAIVE ATTACK.

Model	ASV	Test Pass	Test Fail	Tokens	Cost (\$ USD)
OpenAI: o1	0.02	92	8	69,339	\$3.54
Gemini 2.5 Pro	0.00	100	0	181,436	NA
Deepseek-Chat	0.03	92	8	38,689	NA
Claude Sonnet-4	0.00	100	0	47,022	\$0.53

This attack method is reminiscent of Promptfoo’s methods as it focuses on having the model falsely classify the maliciousness of the message. This jailbreaking method was somewhat effective considering the only thing it did was randomize the location and capitalization of letters within the prompt, however it couldn’t be used to make the LLM follow the user’s instructions. If an attacker wanted to use this method to bypass an AI chat filter it would rarely ever be effective.

In most cases, when faced with a “thought experiment”, the LLM will adhere exactly to its guidelines and will refuse to answer the prompt. Even in cases where the attack is “successful” the output tended to be rather broad and didn’t seem like it would effectively help an attacker in actual attack scenarios. Overall, this attack method is rather ineffective at jailbreaking the model.

The history attack has the most potential to be successful if the attacker follows up correctly. In cases where the attack

is successful, the LLM model would go rather in-depth in how they work and what to look-up to follow up on what it is saying. This format tends to bypass most “I cannot answer that” type of flags, so even in the cases of benign outputs from the LLM, at the very least, they will get a lengthy output.

Consistently there has been a trend of OpenAI LLMs costing significantly more than other models, despite the fact that their tokens total are on par with or less than similar models. Out of all the tests, I’ve come to the conclusion that Claude Sonnet-4 has higher or on par accuracy with other current models, however has a lower cost and token total, making it a more affordable/cost efficient model. OpenAI LLMs have shown a significant vulnerability to prompt injection attacks compared to other models, based on replies, it appears this a result of the LLM not fully “thinking” it

TABLE XIV. RESULTS OF 100 PROMPTS FOR THE “PROMPT INJECTION ATTACK” DATASET ON HUGGINGFACE WHICH CAN BE FOUND AT [XXZ224/PROMPT-INJECTION-ATTACK-DATASET · DATASETS AT HUGGING FACE](#). TESTED UTILIZING PROMPTFOO: LLM EVALUATIONS & RED TEAMING SCANNER. THIS TEST WAS ON THE PROMPT INJECTION’S IGNORE CONTEXT ATTACK.

Model	ASV	Test Pass	Test Fail	Tokens	Cost (\$ USD)
OpenAI: o1	0.08	59	41	88,973	\$4.70
Gemini 2.5 Pro	0.00	100	0	188,773	NA
Deepseek-Chat	0.02	93	7	36,324	NA
Claude Sonnet-4	0.00	100	0	46,100	\$0.51

TABLE XV. RESULTS OF 100 PROMPTS FOR THE “PROMPT INJECTION ATTACK” DATASET ON HUGGINGFACE WHICH CAN BE FOUND AT [XXZ224/PROMPT-INJECTION-ATTACK-DATASET · DATASETS AT HUGGING FACE](#). TESTED UTILIZING PROMPTFOO: LLM EVALUATIONS & RED TEAMING SCANNER. THIS TEST WAS ON THE PROMPT INJECTION’S FAKE COMPLETION ATTACK.

Model	ASV	Test Pass	Test Fail	Tokens	Cost (\$ USD)
OpenAI: o1	0.02	93	7	79,632	\$4.14
Gemini 2.5 Pro	0.01	98	1	187,697	NA
Deepseek-Chat	0.05	88	11	38,516	NA
Claude Sonnet-4	0.00	99	1	47,340	\$0.53

TABLE XVI. RESULTS OF 100 FROM THE “HALLUCINATION” ATTACK METHOD ON FUZZYAI’S ADV_PROMPTS.TXT DATASET. “USES HALLUCINATIONS TO BYPASS RLHF FILTERS”

Model	ASV	Vulnerability Level	Test Pass	Test Fail
GPT 4o	0.02	Very Low Vulnerability	98	2
Deepseek-Chat	0.06	Very Low Vulnerability	94	6
Claude Sonnet-4	0.06	Very Low Vulnerability	94	6

TABLE XVII. RESULTS OF 100 FROM THE “BEST-OF-N JAILBREAKING” ATTACK METHOD ON FUZZYAI’S ADV_PROMPTS.TXT DATASET. “EXPLOITS PROMPT AUGMENTATIONS TO ELICIT HARMFUL RESPONSES FROM AI MODELS ACROSS MODALITIES, ACHIEVING HIGH SUCCESS RATES WITH REPEATED SAMPLING.”

Model	ASV	Vulnerability Level	Test Pass	Test Fail
GPT 4o	0.02	Very Low Vulnerability	98	2

Deepseek-Chat	0.04	Very Low Vulnerability	96	4
Claude Sonnet-4	0.03	Very Low Vulnerability	97	3

TABLE XVIII. RESULTS OF 100 FROM THE “THOUGHT EXPERIMENT” ATTACK METHOD ON FUZZYAI’S ADV_PROMPTS.TXT DATASET. “MODIFIES THE PROMPT BY ADDING A THOUGHT EXPERIMENT-RELATED PREFIX AND PRECAUTIONS IMPLEMENTED SUFFIX.”

Model	ASV	Vulnerability Level	Test Pass	Test Fail
GPT 4o	0.01	Very Low Vulnerability	99	1
Deepseek-Chat	0.01	Very Low Vulnerability	99	1
Claude Sonnet-4	0.02	Very Low Vulnerability	98	2

TABLE XIX. RESULTS OF 100 FROM THE “HISTORY/ACADEMIC FRAMING” ATTACK METHOD ON FUZZYAI’S ADV_PROMPTS.TXT DATASET. “FRAMING SENSITIVE TECHNICAL INFORMATION WITHIN A LEGITIMATE SCHOLARLY OR HISTORICAL RESEARCH CONTEXT TO ENSURE ETHICAL AND LEGAL USE OF KNOWLEDGE, LEADING TO A POTENTIAL JAILBREAK”

Model	ASV	Vulnerability Level	Test Pass	Test Fail
GPT 4o	0.05	Very Low Vulnerability	95	5
Deepseek-Chat	0.04	Very Low Vulnerability	96	4
Claude Sonnet-4	0.04	Very Low Vulnerability	96	4

through in most cases compared to other models. When asked to provide step-by-step reasoning, OpenAI tends to fail in that category due to its overall concise and short-reasoning, despite the fact that it tends to use more tokens than models like Claude that output longer text. Furthermore, it is important to note that Promptfoo considers a fail condition

TABLE XX. RESULTS OF 1000 PROMPTS FOR THE “ARC CHALLENGE” DATASET ON HUGGINGFACE WHICH CAN BE FOUND AT [ALLENAI/ai2_ARC · DATASETS AT HUGGING FACE](#). TESTED UTILIZING PROMPTFOO: LLM EVALUATIONS & RED TEAMING SCANNER.

Model	ASV	Test Pass	Test Fail	Tokens	Cost (\$ USD)
OpenAI: o1	0.03	895	105	613,214	\$31.60
Gemini 2.5 Pro	0.01	958	42	1,759,139	NA
Deepseek-Chat	0.02	950	50	345,244	NA
Claude Sonnet-4	0.01	958	42	467,793	\$5.49

when an output doesn’t meet a singular expectation, OpenAI tends to fail the step-by-step reasoning expectation more than other models, which might skew the results to seeing OpenAI models being significantly inferior. Another interesting condition is that some models monologue for so long that they reach their token limit before they are able to output an answer, in such scenarios, this is considered a fail condition by Promptfoo. Excluding these outliers, it is apparent that most current AI models are within a similar level of cybersecurity regarding prompt injections, however different prompts may lead to a different outcome. For instance, when utilizing the “Prompt Injection” dataset, the prompt asked the model to determine whether or not the message was a prompt injection. In a more authentic scenario where the LLM is unaware of the threat of prompt injection, it appears that LLMs fail to respond accordingly in most cases.

XII. LARGE DATASETS

Due to the potential for LLMs to learn and adapt as they are prompted. It was necessary to see if a larger number of prompts would cause a significant deviation from the usual tests. As a result, it was decided that 1000 prompts from the usual 100 would be run to see what would happen. The ARC-Challenge dataset was chosen to be tested for this due to it having a lack of outliers compared to other datasets. If the results were consistent no matter the size then multiplying the results of the 100 prompt test by 10 should realistically result in a similar output seen in the results from a large test.

The models slightly performed better with their expected token amount and associated cost. That being said, the ASV value and the test passing were less than 1 percent different than the 100 prompt test. As a result, it is reasonable to assume that a smaller dataset can feasibly act in the place of a larger dataset when performing tests using the LLM vulnerability scanner. In fact, the larger dataset caused issues due to Promptfoo not being built to handle more than 100 prompts, which led to some display errors on the GUI side.

XIII. CONCLUSION AND FUTURE WORK

The length of time and CPU usage may vary. It is crucial to consider high traffic during peak hours where higher demand may result in server-side errors or slower response times.

Additionally, costs may vary depending on the persona, how the user adjusts token limits, and the prompt itself. A difference in wording of the persona or hardened prompt may result in a vastly different generated response from the model, so it is crucial for the user to experiment with different personas on limited tests to figure out the best fit.

Future work may involve testing open-source LLMs against the closed-source LLMs tested in this paper. Additionally, we may test different wordings of prompt hardening personas to see how they affect a model's generated response.

ACKNOWLEDGMENT

This work was supported in part by the Commonwealth Cyber Initiative, an investment in the advancement of cyber

R&D, innovation, and workforce development. For more information about CCI, visit cyberinitiative.org.

REFERENCES

- [1] Y. Liu, Y. Jia, R. Geng, J. Jia, and N. Z. Gong, Philadelphia, PA, USA. Formalizing and Benchmarking Prompt Injection Attacks and Defenses. (Aug. 14-16, 2024). Accessed: Sep. 16, 2025. [Online Video]. Available: <https://www.usenix.org/conference/usenixsecurity24/presentation/liu-yupe>
- [2] X. Liu, Z. Yu, Y. Zhang, N. Zhang, and C. Xiao, "Automatic and universal prompt injection attacks against large language models," 2024, arxiv:2403.04957.
- [3] F. Perez and I. Ribeiro, "Ignore previous prompt: Attack techniques for language models," 2022, arxiv:2211.09527.
- [4] S. Lee, J. Kim, and W. Pak, "Mind mapping prompt injection: Visual prompt injection attacks in modern large language models," *Electronics*, vol. 14, no. 10, May 2025, Art. No. 1907, doi: 10.3390/electronics14101907.
- [5] Q. Zhan, R. Fang, H. S. Panchal, and D. Kang, "Adaptive attacks break defenses against indirect prompt injection Attacks on LLM agents," 2025, arxiv:2503.00061.
- [6] K. Hines et al., "Defending against indirect prompt injection attacks with spotlighting," 2024, arxiv:2403.14720.
- [7] X. Suo, "Signed-prompt: A new approach to prevent prompt injection attacks against LLM-integrated applications," 2025, arxiv:2401.07612.
- [8] T. Shi et al., "PromptArmor: Simple yet effective prompt injection defenses," 2025, arxiv:2507.15219.
- [9] S. Chen, J. Piet, C. Sitawarin, and D. Wagner. "StruQ: Defending against Prompt Injection with Structured Queries." *Usenix.org*. Accessed: Sep. 22, 2025. [Online]. Available: <https://www.usenix.org/system/files/usenixsecurity25-chen-sizhe.pdf>
- [10] Y. Jia et al., "A critical evaluation of defenses against prompt injection attacks," 2025, arxiv:2505.18333.
- [11] Y. Liu et al., "Prompt injection attack against LLM-integrated applications," 2024, arxiv:2306.05499.
- [12] M. Alizadeh, Z. Samei, D. Stetsenko, and F. Gilardi, "Simple prompt injection attacks can leak personal data observed by LLM agents during task execution," 2025, arxiv:2506.01055.
- [13] J. Wang, P. Gupta, I. Habernal, and E. Hullermeier, "Is your prompt safe? Investigating prompt injection attacks against open-source LLMs," 2025, arxiv:2505.14368.
- [14] E. Matthew, "Enhancing Security in Large Language Models: A Comprehensive Review of Prompt Injection Attacks and Defenses." *TechRxiv* Accessed: Sep. 22, 2025. [Online]. doi: 10.36227/techrxiv.172954263.32914470/v1 .