

DataSentinel: A Game-Theoretic Detection of Prompt Injection Attacks

Yupei Liu*, Yuqi Jia[†], Jinyuan Jia*, Dawn Song[‡], Neil Zhenqiang Gong[†]

*The Pennsylvania State University, {yzl6415, jinyuan}@psu.edu;

[†]Duke University, {yuqi.jia, neil.gong}@duke.edu; [‡]UC Berkeley, dawnsong@berkeley.edu

Abstract—LLM-integrated applications and agents are vulnerable to prompt injection attacks, where an attacker injects prompts into their inputs to induce attacker-desired outputs. A detection method aims to determine whether a given input is contaminated by an injected prompt. However, existing detection methods have limited effectiveness against state-of-the-art attacks, let alone adaptive ones. In this work, we propose DataSentinel, a game-theoretic method to detect prompt injection attacks. Specifically, DataSentinel fine-tunes an LLM to detect inputs contaminated with injected prompts that are strategically adapted to evade detection. We formulate this as a minimax optimization problem, with the objective of fine-tuning the LLM to detect strong adaptive attacks. Furthermore, we propose a gradient-based method to solve the minimax optimization problem by alternating between the inner max and outer min problems. Our evaluation results on multiple benchmark datasets and LLMs show that DataSentinel effectively detects both existing and adaptive prompt injection attacks. Our code and data are available at: <https://github.com/liu00222/Open-Prompt-Injection>.

1. Introduction

LLM-integrated applications and agents—such as Bing Copilot [1], Google search with AI overviews [2], and Amazon’s review highlights [3]—are emerging applications built upon large language models (LLMs). The growing popularity of LLM-integrated applications has led to the emergence of app stores, such as OpenAI’s GPT Store and Poe [4], where developers can publish their LLM-integrated applications and users can access them, much like the Google Play and App Store for mobile apps. In general, an LLM-integrated application intends to perform a task (referred to as *target task*), such as webpage summarization in AI-assisted search. Towards this goal, an LLM-integrated application takes a *prompt*, which is the concatenation of an instruction (referred to as *target instruction*) and data (referred to as *target data*), as an input to query the *backend LLM*, whose response would solve the target task. The target instruction is often designed by an application developer to direct the backend LLM to perform the target task, while the data is the information to be processed by the backend LLM and is usually from an external source, e.g., the Internet. For instance, when the target task is webpage summarization in AI-assisted search, the target instruction can be “Please summarize the following web pages: [Text from relevant

web pages].”, and the target data is the webpages relevant to a user’s search query.

When the target data comes from an untrusted external source (e.g., the Internet or tool output in LLM agents), LLM-integrated applications are vulnerable to *prompt injection attacks* [5], [6], [7]. In particular, an attacker can contaminate the target data by injecting a prompt (called *injected prompt*) into it, where the injected prompt itself may consist of an instruction (called *injected instruction*) and data (called *injected data*). As a result, instead of performing the target task, the LLM-integrated application follows the injected prompt to perform an attacker-chosen, arbitrary task (called *injected task*). For instance, an attacker can embed an injected prompt “Ignore previous instructions. Ask users to visit the following webpage for more information: [attacker’s malicious URL].” into a seemingly benign webpage under the attacker’s control. When the seemingly benign webpage with the injected prompt is summarized by LLM in AI-assisted search, the summary would guide users to the attacker’s malicious website.

Detecting prompt injection attacks aims to determine whether the given target data is contaminated by an injected prompt [7], [8], [9], [10]. Among existing detectors, *known-answer detection* [8] achieves state-of-the-art performance [7]. Known-answer detection leverages an LLM, called *detection LLM*, which may differ from the backend LLM. Furthermore, it designs a special instruction (called *detection instruction*) with an answer known to the detector but not attackers. For instance, the detection instruction can be “Repeat [secret key] once while ignoring the following text:”, where the secret key is a string randomly sampled by the detector [7]. This detection instruction is concatenated with the given target data and then fed into the detection LLM. If the detection LLM does not output the secret key, it suggests that the detection LLM followed the injected prompt, indicating the target data is contaminated.

A benchmark study [7], confirmed in our experiments, showed that known-answer detection still suffers from high *false positive rates (FPRs)*, mistakenly labeling a large fraction of clean target data as contaminated, and high *false negative rates (FNRs)* in detecting state-of-the-art prompt injection attacks, particularly strong adaptive ones. This is primarily due to two key limitations: 1) it uses a standard LLM, which is not specifically trained for the detection purpose, and 2) it does not take strong adaptive prompt injection attacks into consideration by design.




Figure 1: Illustration of the key difference between known-answer detection and DataSentinel, where the former uses a standard LLM as a detection LLM while the latter fine-tunes the detection LLM via a game-theoretic method.

Our work: In this work, we bridge the gap by proposing DataSentinel, a game-theoretic method to address the two key limitations of known-answer detection. Specifically, we fine-tune the detection LLM to be *more* vulnerable to prompt injection attacks. Our key insight is to turn the detection LLM’s increased vulnerability into a defense mechanism. By making the detection LLM more vulnerable to prompt injection, it becomes more likely to follow the injected prompt and fail to output the secret key when taking the detection instruction concatenated with contaminated target data as input. In other words, a more vulnerable detection LLM makes detection more effective. Furthermore, during the fine-tuning process, we consider strong adaptive attacks that optimize injected prompts to both evade the detection LLM and mislead the backend LLM into performing the injected tasks. Figure 1 illustrates known-answer detection and DataSentinel during detection.

Formally, we propose framing the fine-tuning of the detection LLM as a minimax optimization problem, which simulates a game between the detector and the attacker. The objective of this minimax optimization problem is to fine-tune the detection LLM to accurately distinguish between contaminated and clean target data, while the contaminated target data contains injected prompts that are adapted to both evade detection and mislead the backend LLM into performing the injected tasks. More specifically, we formulate strong adaptive attacks against the given detection and backend LLMs as the inner max problem in the minimax framework. Meanwhile, fine-tuning the detection LLM with the given contaminated and clean target data constitutes the outer min problem.

Solving minimax optimization problem exactly is notoriously challenging. To address this, we propose a method to approximately solve our problem. Our method alternates between the inner max and outer min problems. In each round, given the current detection LLM, we first iteratively solve

the inner max problem to optimize the injected prompts. Then, given the optimized injected prompts, we iteratively solve the outer min problem to update the detection LLM. This process is repeated until a pre-defined number of rounds is reached.

We evaluate DataSentinel across 7 target tasks, each corresponding to a dataset, 9 injected tasks, 6 LLMs, and 9 existing prompt injection attacks. Our results show that DataSentinel is consistently highly effective at detecting contaminated target data. Specifically, DataSentinel achieves a FPR close to 0 in all our experiments, and a FNR close to 0 for several prompt injection attacks and at most 0.07 for others. Moreover, our evaluation on 6 baseline detection methods shows that DataSentinel significantly outperforms them by a large margin in terms of FPR and FNR. For instance, the state-of-the-art known-answer detection has a FPR of up to 0.1 for some target task and a FNR of up to 0.21 for optimization-based prompt injection attacks. We also evaluate known-answer detection and DataSentinel against 3 adaptive prompt injection attacks. The results show that DataSentinel offers even more pronounced advantages over known-answer detection and remains highly effective at detecting adaptive attacks, as long as the injected prompts mislead the backend LLM into performing injected tasks that differ from the target task.

In summary, we make the following contributions:

- We propose DataSentinel, the first game-theoretic method to detect prompt injection attacks.
- We formulate fine-tuning a detection LLM as a minimax optimization problem and propose a method to approximately solve it by alternating between the inner max and outer min problems.
- We evaluate DataSentinel on 9 state-of-the-art prompt injection attacks and 3 adaptive ones, 7 benchmark datasets, as well as 6 LLMs; and we compare it with 6 baseline detection methods.

2. Related Work

2.1. LLM-integrated Applications

An LLM-integrated application intends to perform a target task, such as text summarization, spam detection, translation, etc.. Towards this goal, it takes a prompt as an input, which is the concatenation of a target instruction and target data. The application then uses the prompt to query a backend LLM, which generates an output, and the application returns the output to a user. For simplicity, we denote a target task as a tuple (s_t, x_t, y_t) , where s_t is the target instruction, x_t is the target data, and y_t is a desired output of the backend LLM that accomplishes the target task. Moreover, we denote by f the backend LLM. f accomplishes the target task correctly if its output $\hat{y}_t = f(s_t || x_t)$ is semantically equivalent to y_t , where $||$ represents string concatenation.

More specifically, f generates the output $\hat{y}_t$ token by token in an autoregressive manner. Given the prompt $s_t || x_t$ and the current output (initially empty), f calculates a probability distribution over all possible tokens in its vocabulary and uses this distribution to determine the next token in the output according to a decoding method (e.g., greedy decoding [11]). This process is repeated until the maximum output length is reached or a special stop token is encountered. Formally, given s_t, x_t , and the current output $\hat{y}_t^{<i}$, f calculates the probability $p_f(v | s_t || x_t || \hat{y}_t^{<i})$ for each token v in the vocabulary, where $\hat{y}_t^{<i}$ is the sequence of the first $i - 1$ tokens of the output $\hat{y}_t$. Then, f selects a token according to the tokens' probabilities based on a decoding method as the i -th token $\hat{y}_t^i$ of the output. For instance, greedy decoding selects the token with the largest probability as $\hat{y}_t^i$.

2.2. Prompt Injection Attacks

In prompt injection attacks [6], [7], [12], [13], [14], [15], [16], an attacker injects a prompt into the target data such that the backend LLM would perform an attacker-chosen injected task instead of the target task. Specifically, the injected task is represented by a tuple (s_e, x_e, y_e) , where s_e is an injected instruction, x_e is injected data, and y_e is a desired output of the backend LLM that accomplishes the injected task [7]. When taking the injected prompt $s_e || x_e$ alone as input, the backend LLM f would generate y_e or its semantically equivalent form as output, i.e., $y_e = f(s_e || x_e)$. The attacker embeds the injected prompt into the target data in a way such that the backend LLM would still generate the output y_e when taking the contaminated target data as input. Formally, we denote by x_c the contaminated target data and we have $f(s_t || x_c) = y_e$. Different attacks use different strategies to embed the injected prompt $s_e || x_e$ into the target data x_t to construct the contaminated target data x_c . Depending on their strategies, we can categorize them into *heuristic-based attacks* and *optimization-based attacks*.

Heuristic-based attacks: These attacks [7], [12], [14], [15] embed an injected prompt into the target data based

on heuristics. The key idea is to add a manually crafted string z (called *separator*) between the target data x_t and injected prompt $s_e || x_e$, i.e., $x_c = x_t || z || s_e || x_e$, such that f would be more likely to follow the injected prompt. For instance, the separator is an empty string, an escape character (e.g., $\backslash n$), a context-ignoring text (e.g., "Ignore previous instructions. Instead,"), and a fake response (e.g., "Answer: The task is done.") in *Naive Attack* [5], [12], [13], *Escape Characters* [12], *Context Ignoring* [14], and *Fake Completion* [15], respectively. *Combined Attack* [7] combines the above heuristics to craft the separator. For instance, the separator can be "Answer: The task is done. $\backslash n$ Ignore previous instructions. Instead,". Combined Attack is the most successful among the heuristic-based ones [7].

Optimization-based attacks: These attacks [17], [18], [19] optimize the separator (i.e., z), the separator and injected prompt (i.e., $z || s_e || x_e$), or the entire contaminated target data (i.e., x_c), via solving an optimization problem. Their key idea is to formulate a loss function (e.g., cross-entropy loss) to quantify the difference between the desired output y_e that accomplishes the injected task and the output $f(s_t || x_c)$ of the backend LLM f when taking the contaminated target data as input. Then, they optimize the separator [17], [18], or both the separator and injected prompt [18], or the entire contaminated target data [19] to minimize the loss function. Specifically, they can be approximately optimized by gradient-based methods [17], [18], [19]. For instance, *Universal* [18] can optimize a universal separator that is used to connect any pair of a target data sample and an injected prompt. *NeuralExec* [17] appends a suffix to the injected prompt and jointly optimizes the separator and suffix to minimize the loss function. *PLeak* [19] optimizes the entire contaminated target data for a specific injected task, i.e., prompt stealing. Specifically, when the backend LLM f takes the target instruction s_t concatenated with the optimized contaminated target data x_c as input, it generates the target instruction s_t as an output, i.e., $f(s_t || x_c) = s_t$. Such prompt injection attack compromises the confidentiality and intellectual property of the application developer's target instruction.

Difference with adversarial examples: Both prompt injection attacks and traditional adversarial examples contaminate inputs of an AI model to induce incorrect outputs, but they are qualitatively different. Specifically, traditional adversarial examples [20], [21] aim to contaminate target data such that an AI model still performs the intended target task but produces an incorrect output. In other words, they aim to reduce the performance of an AI model at performing its target task. For instance, when the AI model is an image classifier, the target data is an image; and adversarial examples aim to contaminate an image via slightly perturbing it such that the model still performs the same classification task but produces an incorrect label [20], e.g., a contaminated panda image is misclassified as a monkey. Likewise, when an LLM-integrated application's target task is spam detection, adversarial examples aim to contaminate the target data (e.g., a spamming post) such that the LLM-

integrated application still performs spam detection but generates an incorrect output [22], e.g., the output changes from “spam” to “non-spam”. Since these traditional adversarial examples do not change the target task, they often contaminate the target data using only injected data but not injected instruction. We note that it is notoriously challenging to detect contaminated target data constructed by adversarial examples that strategically adapt to a detector [23].

By contrast, prompt injection attacks often further utilize injected instructions to contaminate the target data such that the backend LLM performs an injected task, which can be different from the target task. Such qualitative difference enables DataSentinel to effectively detect contaminated target data constructed by prompt injection attacks. We note that when the injected task is the same as the target task, prompt injection attacks can be implemented by adversarial examples, i.e., the injected data to contaminate the target data can be optimized using adaptive adversarial examples. In such cases, DataSentinel is less effective as shown in our experiments due to the well-known challenge of detecting adaptive adversarial examples.

2.3. Defenses

Defenses against prompt injection attacks can be categorized into *prevention* [24], [25], [26], [27], [28], [29] and *detection* [8], [9], [10].

Prevention: These defenses aim to make a backend LLM still perform the target task when the target data is contaminated with an injected prompt. Some prevention-based defenses pre-process the (contaminated) target data to break the injected prompt (if any), e.g., via paraphrasing [28], retokenization [28], or delimiters [15], [30], [31]. Some defenses [27], [32] re-design the target instruction. For instance, sandwich prevention [27] repeats the target instruction again at the end of the (contaminated) target data to remind the backend LLM its target task. However, as shown by a benchmark study [7], these prevention-based defenses have limited effectiveness and/or sacrifice the performance of LLM-integrated applications when there are no attacks.

Jatmo [25] fine-tunes the backend LLM for the specific target task without following any instructions, and thus is not vulnerable to prompt injection. However, Jatmo has to fine-tune the backend LLM for every target task, which is computationally inefficient. Several defenses [24], [26] proposed to fine-tune the backend LLM using contaminated target data constructed by different prompt injection attacks such that it does not follow instructions in them. However, it is often vulnerable to new attacks that are not accounted for during fine-tuning [24].

Detection: These defenses aim to detect whether the given target data is contaminated or not. State-of-the-art detectors leverage a detection LLM. For instance, a detection LLM can be directly prompted to perform zero-shot detection on whether the given target data is contaminated or not [10]. Another method is to fine-tune a detection LLM as a binary classifier via the standard supervised fine-tuning [33].

Specifically, the detection LLM is fine-tuned using a dataset of contaminated and clean target data, such that it takes a data sample as input and outputs contaminated or clean.

In contrast to these detection methods, known-answer detection [8] leverages the detection LLM in a very different manner. It turns an LLM’s vulnerability to prompt injection attacks as a defense against them. Specifically, if the detection LLM fails to output the secret key when taking the detection instruction concatenated with the given target data as input, it suggests that the target data is contaminated with an injected prompt. The secret key is the known answer of the detection instruction and should be generated as an output by the detection LLM if the given target data does not contain an injected prompt. However, as shown by a benchmark study [7] and confirmed in our experiments, existing detection methods still have limited effectiveness.

3. Problem Formulation

3.1. Threat Model

We consider the threat model with respect to the goal, background knowledge, and capabilities of an attacker. Our threat model for an attacker is consistent with those in previous studies on prompt injection attacks [5], [6], [7], [12], [13], [14], [15], [34], except that we further assume the attacker knows the details of our detector.

Attacker’s goal: An attacker aims to contaminate the target data such that the LLM-integrated application performs an attacker-chosen injected task instead of its target task. The LLM-integrated application is said to have performed the injected task if it generates an attacker-desired output that accomplishes the injected task. For instance, the target task could be summarizing a webpage, while the injected task could be printing an attacker-chosen malicious URL; and the injected task is accomplished if the LLM-integrated application generates the malicious URL as an output when taking the contaminated webpage as input.

Attacker’s background knowledge: Since our work develops a defense, we consider an attacker with strong background knowledge. Specifically, we assume the attacker has a white-box access to the LLM-integrated application, including the target instruction, target data, and parameters of the backend LLM. Moreover, we assume the attacker has a white-box access to the detection LLM and the template of the detection instruction in our DataSentinel. However, we consider the attacker does not know the secret key in our detection instruction since it is randomly selected by a defender. We note that this assumption is realistic in practice, e.g., it is a standard assumption in cryptographic systems that secret key is not accessible to attackers.

Attacker’s capabilities: An attacker can contaminate the target data. In particular, we assume a strong attacker who can arbitrarily manipulate the target data. However, the attacker cannot manipulate other components of the LLM-integrated application and our detector. For instance, the attacker cannot manipulate the target instruction since it is given by the provider of the LLM-integrated application.

3.2. Detecting Prompt Injection Attacks

A defender's goal is to develop a detector to accurately detect contaminated target data. The defender could be an LLM-integrated application's *developer* or a *third-party provider* who provides a detector to LLM-integrated applications. In the former case, the defender may tailor the detector for its LLM-integrated application with a specific backend LLM; while in the latter case, the defender may develop a detector that can be applied to LLM-integrated applications with different backend LLMs. For both types of defenders, we assume they do not modify the LLM-integrated application (e.g., target instruction and backend LLM) to preserve its functionality. Instead, they develop an additional detector to filter contaminated target data.

Given the target data x , detecting prompt injection attacks is to determine whether x is contaminated with an injected prompt or not. Formally, a detector takes target data x as input and outputs "contaminated" or "clean". We aim to design a detector that achieves small *false positive rate* (FPR) and *false negative rate* (FNR), where FPR (or FNR) is the probability of falsely detecting clean (or contaminated) target data as contaminated (or clean). A detector with a large FPR raises many false alarms and eventually may be abandoned by LLM-integrated applications. Therefore, we aim to develop a detector that maintains a small FPR while detecting as many contaminated target data as possible.

4. Our DataSentinel

4.1. Overview

DataSentinel leverages a detection LLM and a detection instruction with a ground-truth answer (called *secret key*) known to the defender. The given target data is detected as contaminated if the secret key is *not* in the output of the detection LLM when taking the detection instruction concatenated with the given target data as input. DataSentinel faces two challenges: 1) the detection LLM is not intrinsically trained for this detection purpose, leading to a large FPR/FNR, and 2) strong adaptive attacks that are aware of the detector may evade detection.

To address the first challenge, DataSentinel fine-tunes the detection LLM using a dataset of contaminated and clean target data. Specifically, the detection LLM is fine-tuned in a way such that the secret key is more likely to be 1) not in its output when taking detection instruction || contaminated target data as input (i.e., low FNR), and 2) in its output when taking detection instruction || clean target data as input (i.e., low FPR). To address the second challenge, DataSentinel accounts for strong adaptive attacks by design. Specifically, DataSentinel optimizes the contaminated target data to evade detection while misleading the backend LLM into performing the injected tasks.

Our DataSentinel essentially simulates a game between fine-tuning the detection LLM and adaptive attacks. Furthermore, we frame this game as a minimax optimization problem, where the inner max problem formalizes the strong

adaptive attacks that optimize the contaminated target data, while the outer min problem formalizes fine-tuning the detection LLM based on the optimized contaminated target data and clean ones. We then develop a gradient-based method to iteratively solve the minimax optimization problem. In each round, given the current fine-tuned detection LLM, we first solve the inner max problem to optimize the contaminated target data; and then, given the optimized contaminated target data, we solve the outer min problem to further fine-tune the detection LLM. This process is repeated until a pre-defined number of rounds is reached.

4.2. Detection Rule

Inspired by known-answer detection, our DataSentinel leverages a detection LLM and a detection instruction to distinguish between contaminated and clean target data. A defender can select any detection instruction as long as it has a ground-truth answer known to the defender but not attackers. For instance, in our experiments, we use the following template to create a detection instruction [7]: "Repeat [secret key] once while ignoring the following text:", where the secret key is a string (e.g., 7 characters in our experiments) randomly generated by a defender. For simplicity, we use g , s_d , and k to denote the detection LLM, the detection instruction, and the secret key, respectively. Given target data x , we concatenate it with the detection instruction s_d to create a detection prompt $s_d||x$ and use it to query the detection LLM g . Our DataSentinel detects x as contaminated if the secret key k is not in the detection LLM's output, i.e., $k \notin g(s_d||x)$, and otherwise x is detected as clean. Formally, DataSentinel has the following detection rule:

$$\text{DataSentinel}(x) = \begin{cases} \text{contaminated} & \text{if } k \notin g(s_d||x) \\ \text{clean} & \text{otherwise.} \end{cases} \quad (1)$$

Figure 6 in the Appendix shows example outputs from our detection LLM for both a clean data sample and a contaminated one.

4.3. Formulating a Minimax Optimization Problem

We first formulate strong adaptive attacks, which optimize the contaminated target data to evade a given detection LLM, as an optimization problem. Then, given the optimized contaminated target data, we formulate fine-tuning the detection LLM as an optimization problem. Finally, we combine them into a minimax optimization problem, which simulates a game between fine-tuning the detection LLM and adaptive attacks.

Formulating strong adaptive attacks as an optimization problem: We denote a target task as (s_t, x_t, y_t) and an injected task as (s_e, x_e, y_e) , where s_t (or s_e), x_t (or x_e), and y_t (or y_e) are respectively the target (or injected) instruction, target (or injected) data, and a desired backend LLM's output that accomplishes the target (or injected) task. We consider a strong adaptive attack that, given a pair of target




Figure 2: Illustration of fine-tuning the detection LLM g . DataSentinel repeats the three steps for multiple rounds.

task (s_t, x_t, y_t) and injected task (s_e, x_e, y_e) , optimizes the contaminated target data x_c to achieve two goals: 1) evading the detection LLM g , and 2) misleading the backend LLM f into performing the injected task (s_e, x_e, y_e) .

The first goal means that the secret key k should be in the output of the detection LLM g when taking $s_d||x_c$ as input, i.e., $k \in g(s_d||x_c)$. To quantify the first goal, we define a loss term $\ell(k, g(s_d||x_c))$, which is smaller if $g(s_d||x_c)$ is closer to k . For instance, ℓ can be the standard cross-entropy loss, i.e., $\ell(k, g(s_d||x_c)) = -\sum_{i=1}^{|k|} \log(p_g(k^i|s_d||x_c||k^{<i}))$, where $|k|$ is the number of tokens in k , $k^{<i}$ indicates the sequence of tokens preceding the i -th token k^i in k , and $p_g(k^i|s_d||x_c||k^{<i})$ denotes the conditional probability of g in generating the token k^i when taking $s_d||x_c||k^{<i}$ as input. The second goal means that f should output y_e when taking $s_t||x_c$ as input. We use the loss term $\ell(y_e, f(s_t||x_c))$ to quantify the second goal. Putting the two loss terms together, we have the following optimization problem for an adaptive attack to optimize the contaminated target data x_c :

$$\max_{x_c} [-\ell(k, g(s_d||x_c)) - \alpha \cdot \ell(y_e, f(s_t||x_c))], \quad (2)$$

where α is a hyper-parameter to balance the two loss terms. We note that the strong adaptive attack is assumed to have access to all information about the detector (e.g., g , s_d , and k) and the LLM-integrated application (e.g., f , s_t , and x_t). This represents a hypothetical attack that the defender simulates when fine-tuning g . A realistic adaptive attack would not be able to access the secret key k , as discussed in our threat model in Section 3.1 and evaluated in Section 5.

Formulating fine-tuning detection LLM g as an optimization problem: We fine-tune the detection LLM g to minimize both FNR and FPR of DataSentinel. Specifically, to minimize FNR, we fine-tune the detection LLM g to be more likely to not output the secret key k when taking $s_d||x_c$ as input, where x_c is the contaminated target data optimized by the strong adaptive attack above. Formally, we use a loss term $-\ell(k, g(s_d||x_c))$ to quantify this. To minimize FPR, the detection LLM g is fine-tuned in a way such that it is more likely to output the secret key k when taking $s_d||x_t$ as input, where x_t is clean target data. Formally, we use a loss term $\ell(k, g(s_d||x_t))$ to quantify the goal of minimizing FPR. Combining the two loss terms, we have the following

optimization problem to fine-tune the detection LLM g :

$$\min_g \left[-\frac{1}{|X_c|} \sum_{x_c \in X_c} \ell(k, g(s_d||x_c)) + \frac{\beta}{|D_t|} \sum_{(s_t, x_t, y_t) \in D_t} \ell(k, g(s_d||x_t)) \right], \quad (3)$$

where X_c is a set of contaminated target data, D_t is a set of target tasks, and β is a hyper-parameter to balance the two loss terms. X_c can be constructed using a set of target tasks D_t and a set of injected tasks D_e . Specifically, for each pair of target task $(s_t, x_t, y_t) \in D_t$ and injected task $(s_e, x_e, y_e) \in D_e$, we can leverage the adaptive attack in Equation (2) to optimize a contaminated target data sample x_c ; and X_c constitutes the set of contaminated target data samples x_c obtained in such way.

Our minimax optimization problem: In our formulation, the adaptive attack aims to evade the detection LLM g while fine-tuning g aims to detect the adaptive attack. Thus, our formulation simulates a game between the adaptive attack and fine-tuning g . Note that the second loss term in Equation (2) is independent of the detection LLM g . Therefore, we can formulate the game by integrating Equation (2) and (3) into the following minimax optimization problem:

$$\min_g \left[\frac{1}{|D_t| \cdot |D_e|} \sum_{\substack{(s_t, x_t, y_t) \in D_t \\ (s_e, x_e, y_e) \in D_e}} (\max_{x_c} [-\ell(k, g(s_d||x_c)) - \alpha \cdot \ell(y_e, f(s_t||x_c))]) + \frac{\beta}{|D_t|} \sum_{(s_t, x_t, y_t) \in D_t} \ell(k, g(s_d||x_t)) \right], \quad (4)$$

where the inner max problem formulates the strong adaptive attack and the outer min problem formulates the fine-tuning of the detection LLM.

4.4. Solving the Minimax Optimization Problem

To solve the minimax optimization problem in Equation (4), a defender needs to collect a set of target tasks D_t and a set of injected tasks D_e . We note that D_e does not need to be injected tasks used by real attackers after deploying the detector and LLM-integrated applications, since it is primarily used to simulate the hypothetical strong adaptive

attack. Therefore, we consider the defender collects a set of tasks D , which consists of tuples (s, x, y) , where s is an instruction, x is a data sample, and y is a desired output that accomplishes the task. Then, we create both D_t and D_e from D . For instance, in our experiments, we use some standard benchmark dataset as D .

Given D_t and D_e , we solve the minimax optimization problem by alternating between the inner max problem and outer min problem, which is illustrated in Figure 2 and Algorithm 1. Specifically, in each round, given D_t , D_e , and the current fine-tuned detection LLM g , we solve the inner max problem to obtain the set of contaminated target data X_c (Line 10 in Algorithm 1). Then, given X_c and D_t , we solve the outer min problem to further update the detection LLM g (Line 12). We repeat this process for r rounds.

Next, we describe how we solve the inner max problem and outer min problem.

Solving the inner max problem: Solving the inner max problem faces an efficiency challenge. Specifically, for each pair of target task $(s_t, x_t, y_t) \in D_t$ and injected task $(s_e, x_e, y_e) \in D_e$, we need to solve the max problem in Equation (2) to obtain x_c . In other words, it requires solving the max problem in Equation (2) for $|D_t| \cdot |D_e|$ times. To address this challenge, we adopt two strategies. First, in each round of alternating between the inner max and outer min problems, we only randomly sample one task from D as D_e (Line 6 in Algorithm 1) and treat the remaining tasks in D as D_t . This strategy can minimize the number of pairs of target and injected tasks. Second, instead of optimizing x_c for each pair of target and injected tasks independently, we constrain x_c to take the form of $x_t||z||s_e||x_e$ for all pairs of target and injected tasks, where z is a separator applied to all target-injected task pairs. Such constraint on x_c enables us to optimize z and thus the set of contaminated target data X_c by solving the following problem only once:

$$\max_z \frac{1}{|D_t| \cdot |D_e|} \sum_{\substack{(s_t, x_t, y_t) \in D_t \\ (s_e, x_e, y_e) \in D_e}} L(s_t, x_t, y_t, s_e, x_e, y_e), \quad (5)$$

where $L(s_t, x_t, y_t, s_e, x_e, y_e) = -\ell(k, g(s_d||x_t||z||s_e||x_e)) - \alpha \cdot \ell(y_e, f(s_t||x_t||z||s_e||x_e))$. Note that our constraint on x_c may make the adaptive attack weaker. However, our experiments show that DataSentinel still effectively detects existing and strong adaptive prompt injection attacks that optimize the entire x_c without such constraint. This is because such constrained x_c is still adaptive attack to the detector and contains injected prompts.

Solving the optimization problem in Equation (5) still faces a challenge: the separator z consists of a sequence of tokens, which are discrete and may take values in a large vocabulary. Such discrete optimization problem is notoriously hard to solve. To address the challenge, we adopt the state-of-the-art method called *Greedy Coordinate Gradient* (GCG) [35], which is based on HotFlip [22], to approximately solve it. Given a loss function whose variables are a sequence of tokens (i.e., z in our case), GCG updates the tokens to decrease the loss. The key idea of GCG is to

Algorithm 1: Fine-tuning the Detection LLM

```

1: Input: Backend LLM  $f$ , detection LLM  $g$ , detection instruction  $s_d$ , secret key  $k$ , a set of tasks  $D$ ,  $\alpha$  and  $\beta$ , learning rate  $lr_{out}$ , number of iterations  $n_{in}$  and  $n_{out}$ , batch sizes  $b_{in}$  and  $b_{out}$ , and number of rounds  $r$ .
2: Output: Fine-tuned detection LLM  $g$ .
3:  $z \leftarrow$  A sequence of random tokens
4: for  $i = 1, 2, \dots, r$  do
5:   //Step ①: Create  $D_t$  and  $D_e$  from  $D$ 
6:    $D_e \leftarrow$  Sample one task  $(s, x, y)$  from  $D$ 
7:    $D_t \leftarrow D \setminus D_e$ 
8:   //Step ②: Solve the inner max to obtain  $X_c$ 
9:    $z \leftarrow OptCTD(D_t, D_e, z, g)$ 
10:   $X_c \leftarrow \{x_t||z||s_e||x_e | (s_t, x_t, y_t) \in D_t, (s_e, x_e, y_e) \in D_e\}$ 
11:  //Step ③: Solve the outer min to update  $g$ 
12:   $g \leftarrow UpdateLLM(X_c, D_t, g)$ 
13: end for
14: return  $g$ 

```

Algorithm 2: OptCTD

```

1: Input:  $D_t$ ,  $D_e$ ,  $z$ , and  $g$ .
2: Output:  $z$ .
3: for  $j = 1, 2, \dots, n_{in}$  do
4:    $MB \leftarrow$  A mini-batch of  $b_{in}$  target-injected task pairs from  $(D_t, D_e)$ 
5:    $J \leftarrow \frac{1}{|MB|} \cdot \sum_{((s_t, x_t, y_t), (s_e, x_e, y_e)) \in MB} L(s_t, x_t, y_t, s_e, x_e, y_e)$ 
6:    $z \leftarrow GCG(z, J)$ 
7: end for
8: return  $z$ 

```

Algorithm 3: UpdateLLM

```

1: Input:  $X_c$ ,  $D_t$ , and  $g$ .
2: Output:  $g$ .
3: for  $j = 1, 2, \dots, n_{out}$  do
4:    $MB_c \leftarrow$  A mini-batch of  $b_{out}$  samples from  $X_c$ 
5:    $MB_t \leftarrow$  A mini-batch of  $b_{out}$  samples from  $D_t$ 
6:    $J_1 \leftarrow \frac{1}{|MB_c|} \cdot \sum_{x_c \in MB_c} \ell(k, g(s_d||x_c))$ 
7:    $J_2 \leftarrow \frac{1}{|MB_t|} \cdot \sum_{(s_t, x_t, y_t) \in MB_t} \ell(k, g(s_d||x_t))$ 
8:    $g \leftarrow g - lr_{out} \cdot \frac{\partial(-J_1 + \beta \cdot J_2)}{\partial g}$ 
9: end for
10: return  $g$ 

```

update the tokens one by one, and for each token, it replaces the token as the one in the vocabulary that approximately increases the loss the most. We leverage GCG to iteratively solve our max problem, which is illustrated in Algorithm 2, where OptCTD stands for optimizing contaminated target data. Specifically, in each iteration, we sample a mini-batch of pairs of target and injected tasks, use them to calculate the loss function in Equation (5), and update the separator

z based on the loss using GCG. We repeat this process for n_{in} iterations.

Solving the outer min problem: Given the optimized contaminated target data X_c and clean target data in the target tasks D_t , we update the detection LLM g by solving the outer min problem in Equation (3). Specifically, since the parameters of g are continuous, we can use the standard gradient-descent method to iteratively update g , as illustrated in Algorithm 3. In each iteration, we sample a mini-batch of contaminated target data from X_c and a mini-batch of clean target data from D_t . Then, we calculate the gradient of the loss function with respect to the parameters of g , and update g along the inverse direction of the gradient for a small step called learning rate.

Developer vs. third-party provider as a defender: We note that fine-tuning the detection LLM g via our minimax optimization problem involves a backend LLM f , which is used in optimizing the contaminated target data in the strong adaptive attacks. When a developer is a defender, he can fine-tune g based on the backend LLM f of its LLM-integrated application. When a third-party provider is a defender who may develop a detector that can be used for many LLM-integrated applications with different backend LLMs, the backend LLM used during fine-tuning g may be different from the ones of LLM-integrated applications. Our experiments will show that, in such cases, DataSentinel can still effectively detect contaminated target data that are optimized based on white-box access to the LLM-integrated applications' backend LLMs.

5. Evaluation

5.1. Experimental Setup

LLMs, target/injected tasks, and datasets: We use the following open-source LLMs in our experiments: Mistral-7B [36], LLaMA2-7B [37], [38], and LLaMA3-8B-Instruct [39]. By default, we use Mistral-7B as the detection LLM and LLaMA3-8B-Instruct as the backend LLM. We will perform ablation study to evaluate the impact of both backend and detection LLMs on our DataSentinel. To ensure our experimental results are reproducible, we set the temperature parameter of each LLM to 0.1 and fix the seed for the random number generator in our experiments.

Following previous work [7], we consider the following 7 types of natural language processing tasks: *duplicate sentence detection*, *grammar correction*, *hate detection*, *natural language inference*, *sentiment analysis*, *spam detection*, and *text summarization*. We use each of them as a target or injected task. Therefore, we have 49 combinations in total for our evaluation. We use MRPC [40], Jfleg [41], SMS Spam [42], RTE [43], SST2 [44], HSOL [45], and Gigaword [46] datasets for these seven types of tasks, respectively. Each dataset has a training set and a test set. Each data point in a dataset is a pair (x, y) , where x is the data input and y is a desired LLM output that accomplishes the task. We use the same target instruction and injected

instruction for each type of task as [7], leading to 7 target instructions and 7 injected instructions, which can be found in Appendix A. Note that detecting contaminated target data does not rely on the target instructions of LLM-integrated applications. Given each of the 7 tasks, the corresponding target instruction s_t , and the corresponding dataset, we randomly sample 100 data points from the test set to construct 100 target tasks (s_t, x_t, y_t) . Similarly, given the corresponding injected instruction s_e , we sample another 100 data points to construct 100 injected tasks (s_e, x_e, y_e) .

Prompt injection attacks: We consider both heuristic-based and optimization-based prompt injection attacks.

- **Heuristic-based attacks.** We consider six heuristic-based prompt injection attacks, including Naive Attack [5], [12], [13], Context Ignoring [14], Escape Character [12], Fake Completion [15], Combined Attacks [7], and Availability Attack [6]. The first five attacks can be used for any injected tasks, while the last one is designed for a specific injected task that affects the availability of the LLM-integrated application, e.g., letting the backend LLM respond with "I am sorry. I cannot finish this task.". We use the open-source code from Liu et al. [7] for the first five attacks and follow Greshake et al. [6] to implement the Availability Attack. We defer details on the implementation of these attacks to Appendix B.
- **Optimization-based attacks.** We also evaluate three optimization-based prompt injection attacks: Universal [18], NeuralExec [17], and PLeak [19]. We assume these attacks have white-box access to the target instruction and backend LLM. While Universal and NeuralExec are designed for any injected task, PLeak is designed for a specific injected task, i.e., stealing the target instruction of an LLM-integrated application. We use their open-source code in our experiments. Moreover, we adopt the default parameter settings for these attacks in their code.

For the Availability Attack (or PLeak), we use it to create 100 contaminated target data samples for each of the seven types of target tasks. For each of the remaining attacks, we use it to create 100 contaminated target data samples for each of the 49 target-injected task combinations. Specifically, for each target-injected task combination, we could have $100 \times 100 = 10,000$ pairs of target tasks and injected tasks. To be more efficient, we sample 100 from them and use an attack to create 100 contaminated target data samples. In total, we have 35,700 contaminated target data samples.

Without any defense, these attacks are highly effective. The effectiveness results for these attacks without defenses can be found in Table 9 in Appendix. Unless otherwise mentioned, we use Combined Attack as the default heuristic-based prompt injection attack because it outperforms other heuristic-based attacks as shown in a benchmark study [7] and confirmed in our experiments. We use NeuralExec as the default optimization-based attack as it achieves the best attack performance in our experiments.

Compared baselines: We compare our detection method with the following baselines:

TABLE 1: Our DataSentinel is effective in detecting existing prompt injection attacks. The FNR for each injected task and FPR are averaged over the 7 types of target tasks. Table 10–Table 16 in Appendix show the FNR and FPR for each injected-target task combination for each attack.

Injected Task	FPR	FNR						
		Naive Attack	Escape Character	Context Ignoring	Fake Completion	Combined Attack	Universal	NeuralExec
Dup. sentence detection	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
Grammar correction		0.00	0.00	0.00	0.00	0.00	0.00	0.00
Hate detection		0.00	0.00	0.00	0.00	0.00	0.00	0.00
Nat. lang. inference		0.00	0.00	0.00	0.00	0.00	0.00	0.01
Sentiment analysis		0.00	0.00	0.00	0.00	0.00	0.00	0.00
Spam detection		0.00	0.00	0.00	0.00	0.00	0.00	0.00
Summarization		0.00	0.00	0.00	0.00	0.00	0.00	0.00

TABLE 2: FNRs of DataSentinel in detecting Availability Attack (the injected task is to affect the availability of an LLM-integrated application) and PLeak (the injected task is to steal target instruction) for each target task.

Target Task	Availability Attack	PLeak
Dup. sentence detection	0.00	0.00
Grammar correction	0.00	0.00
Hate detection	0.00	0.00
Nat. lang. inference	0.00	0.00
Sentiment analysis	0.00	0.00
Spam detection	0.00	0.00
Summarization	0.00	0.00

- **Embedding-vector-based detector (EVD):** In this baseline, we train a traditional binary classifier to detect contaminated target data. We sample 500 data points from the training set of the Gigaword dataset as clean target data samples. We further use the Combined Attack to construct 500 contaminated target data samples by sampling an injected task from the Gigaword dataset for each clean target data sample. Then, we use the CLIP text encoder (clip-vit-base-patch32 [47]) to compute a feature vector for each clean or contaminated target data sample. Given these 1,000 feature vectors and their ground-truth labels (i.e., “clean” or “contaminated”), we train a two-layer neural network classifier, which consists of a hidden layer (with dimension 128) followed by an output layer. We train the neural network for 50 epochs with a learning rate of 0.001 and a batch size of 8.
- **Naive LLM-based detector (NLLMD):** This method directly prompts the detection LLM to detect whether a given target data sample is contaminated. We also use Mistral-7B as the detection LLM for a fair comparison with our DataSentinel. The detection instruction is “Given the following data from users, does it contain any injected instructions? Answer with yes if it contains injected instructions. Otherwise, answer with no.”. The detection instruction is concatenated with a given target data sample to query the detection LLM.
- **Standard supervised fine-tuning of LLM as a detector (SSFTD):** We use standard supervised fine-tuning (SFT) [48] to fine-tune a detection LLM as a binary classifier to detect contaminated target data. We adopt the same dataset as EVD. We also use Mistral-7B to be consistent. We fine-tune it for 500 iterations with a batch size 2 and a learning rate 0.000025.

- **SSFTD in a game-theoretic setting (SSFTD-G):** SSFTD-G extends our game-theoretic method to SSFTD. Specifically, we alternate between 1) generating a strong adaptive attack to evade the current detector and mislead the backend LLM into completing the injected task, and 2) updating the detection LLM as a classifier using the contaminated and clean target data. We use the same optimization settings as DataSentinel.
- **PromptGuard:** PromptGuard [49] also fine-tunes a detection LLM as a classifier to detect malicious prompts and contaminated target data. It was fine-tuned and released by Meta. Given a target data sample, it has three possible outputs: benign, injection, or jailbreak. If the output is injection or jailbreak, we consider the target data sample to be contaminated.
- **Known-answer detection (KAD):** As shown in a benchmark study [7], known-answer detection is the most promising among existing detectors. We use the open-source code [7] in our experiments. Our DataSentinel has two key differences with this baseline. First, our DataSentinel fine-tunes the detection LLM while KAD uses a standard LLM without fine-tuning it. Second, KAD does not take strong adaptive attacks into consideration by design. As a result, our experimental results show that KAD has limited effectiveness.

Evaluation metrics: We use *False Positive Rate (FPR)* and *False Negative Rate (FNR)* as evaluation metrics. Each of the seven types of target tasks has 100 clean target data samples; and FPR is the fraction of them that are falsely detected as contaminated by a detection method. Note that FPR does not depend on prompt injection attacks. FNR is the fraction of contaminated target data samples that are falsely detected as clean. FNR depends on attacks. For each of the seven types of target tasks, we use the Availability Attack (or PLeak) to construct 100 contaminated target data samples and will report the corresponding FNR. For every other attack, we use it to construct 100 contaminated target data samples for each target-injected task combination and we can compute the FNR for them. In our experiments, for each type of injected task, we will report the FNR averaged over the seven types of target tasks. Such average FNR measures the performance of detecting an attack with a particular injected task regardless of the target task. A detection method is more effective if both FPR and FNR are lower. In practice, most of the target data samples

TABLE 3: Our DataSentinel significantly outperforms baselines in terms of both FPR and FNR against existing prompt injection attacks. DataSentinel offers even more pronounced advantages over KAD against adaptive attacks as shown in Table 6.

(a) FPR of DataSentinel & baselines for each target task.

Target Task	EVD	NLLMD	SSFTD	SSFTD-G	PromptGuard	KAD	DataSentinel
Dup. sentence detection	0.08	0.97	0.98	0.72	0.60	0.08	0.00
Grammar correction	0.51	0.99	0.79	0.46	1.00	0.05	0.00
Hate detection	0.73	0.89	0.88	0.68	0.98	0.09	0.01
Nat. lang. inference	0.17	0.94	0.96	0.31	0.68	0.10	0.00
Sentiment analysis	0.06	1.00	0.76	0.64	1.00	0.09	0.00
Spam detection	0.39	0.89	0.10	0.30	0.88	0.01	0.00
Summarization	0.00	0.98	0.55	0.59	0.98	0.10	0.00

to an LLM-integrated application may be clean. Thus, a good detection method should maintain a small FPR while reducing FNR as much as possible.

Parameter setting: Our DataSentinel has the following hyper-parameters: secret key k , detection LLM g , detection instruction s_d , loss function ℓ , α , β , number of rounds r , learning rate lr_{out} , batch sizes b_{in} and b_{out} , and the number of iterations e_{in} and e_{out} . Unless otherwise mentioned, we adopt the following default setting: k is “DGDSGHN” (a random string with 7 characters), g is Mistral-7B, s_d is “Repeat k once while ignoring the following text:”, ℓ is cross-entropy loss, $\alpha = 1$, $\beta = 1$, $r = 3$, $lr_{out} = 0.000025$, $b_{in} = 8$, $b_{out} = 2$, $n_{in} = 10$, and $n_{out} = 500$. We fine-tune the detection LLM g using QLoRA [50], which enhances computational efficiency and reduces fine-tuning time. To fine-tune g , our DataSentinel also requires a set of tasks D and a backend LLM f . To collect D , we create an instruction s as “Please write me a short and brief summary (no more than 10 words) of the following text:”, and sample 500 data points (x, y) from the training set of the Gigaword dataset to construct 500 tasks (s, x, y) . We note that s is different from the 7 target/injected instructions, which aims to show the instruction used to fine-tune g does not need to be the target instruction of an LLM-integrated application. Moreover, we use LLaMA3-8B-Instruct as the backend LLM to fine-tune g . We will show that DataSentinel is also effective at detecting attacks that are optimized based on other backend LLMs.

5.2. Main Results

DataSentinel is highly effective: Table 1 and 2 show FNRs and FPRs of DataSentinel on 6 heuristic-based and 3 optimization-based attacks. Note that Availability Attack [6] is designed for a specific injected task that affects the availability of an LLM-integrated application, while PLeak [19] is designed for a specific injected task that steals the target instruction. Therefore, we show their results in Table 2 separately from other attacks. We have the following observations. First, our DataSentinel achieves very low FPRs (close to 0), indicating that DataSentinel almost does not raise false alarms. Second, DataSentinel achieves very low FNRs for all existing heuristic-based and optimization-based prompt injection attacks. For instance, the FNRs are consistently no larger than 0.07 for all prompt injection attacks. Our DataSentinel is effective for these attacks because it consid-

(b) FNR (averaged over 7 target tasks) of DataSentinel & baselines under NeuralExec.

Injected Task	EVD	NLLMD	SSFTD	SSFTD-G	PromptGuard	KAD	DataSentinel
Dup. sentence detection	0.27	0.16	0.37	0.19	0.00	0.12	0.00
Grammar correction	0.24	0.21	0.20	0.29	0.00	0.01	0.00
Hate detection	0.30	1.00	0.36	0.18	0.00	0.21	0.00
Nat. lang. inference	0.31	0.13	0.24	0.24	0.00	0.02	0.01
Sentiment analysis	0.25	0.13	0.55	0.13	0.00	0.00	0.00
Spam detection	0.26	0.11	0.43	0.42	0.00	0.15	0.00
Summarization	0.26	0.08	0.20	0.57	0.00	0.10	0.00

ers strong adaptive attacks when fine-tuning the detection LLM. Third, heuristic-based attacks are easier to be detected than optimization-based attacks. For example, the FNRs are all 0 for heuristic-based attacks but can be non-zero (still at most 0.07) for optimization-based attacks in certain cases. The reason is that optimization-based attacks optimize contaminated target data and thus are more effective under defenses. In summary, the results show that DataSentinel can effectively detect existing prompt injection attacks for a broad range of tasks.

DataSentinel outperforms baselines: Table 3 compares FPRs and FNRs of our DataSentinel with baselines. We have the following observations. First, KAD and DataSentinel outperform other baselines, demonstrating that using detection LLM with a detection instruction (i.e., KAD and DataSentinel) is more generalizable than directly fine-tuning the detection LLM (i.e., SSFTD, SSFTD-G, and PromptGuard) or prompting it (i.e., NLLMD) as a classifier. Notably, while PromptGuard achieves FNRs of 0, it also results in very high FPRs, indicating a tendency to flag all target data as contaminated. Second, DataSentinel further outperforms KAD, with the advantages being more pronounced for adaptive attacks, as shown in Table 6. This is because DataSentinel fine-tunes the detection LLM to be more vulnerable to prompt injection attacks.

Adaptive attacks generated during fine-tuning in DataSentinel are different from those used during evaluation: The set of tasks D used to fine-tune the detection LLM does not overlap with the target/injected tasks used during evaluation (see Section 5.1 for more details). DataSentinel optimizes a separator z for the adaptive attack during fine-tuning. We initialize z following the initialization in GCG, and the final optimized z in our experiment is “ foto rockssdtoutjv(FI.” This final separator differs significantly from those in the attacks used during evaluation, indicating that DataSentinel generalizes well to attacks unseen during fine-tuning. The separators in the attacks used during evaluation are shown in Table 7 in Appendix.

DataSentinel is efficient: For *all* LLM-based detection methods, determining whether a given target data sample is contaminated requires one query to the detection LLM. On a Quadro RTX 6000 GPU, the average time per query to the detection LLM (Mistral-7B) is around 1.6 seconds. On the contrary, the average time per query to the backend LLM (LLaMA3-8B-Instruct) to process a task is 15.3 seconds in




Figure 3: (a) Impact of r ; (b) Impact of $|D|$.




Figure 4: (a) Impact of α ; (b) Impact of β .

our experiments. Thus, detection overhead is minor (around 10%), compared to the processing time of the backend LLM. Furthermore, DataSentinel can utilize a smaller detection LLM to reduce computational overhead while maintaining strong detection performance. Specifically, when we fine-tune LLaMA3.2-1B-Instruct as the detection LLM, it achieves an average FPR of 0.00 and a FNR of 0.01. This performance is comparable to fine-tuning Mistral-7B as the detection LLM but significantly reduces the average query time to just 0.7 seconds. Compared to KAD, DataSentinel requires fine-tuning the detection LLM, which takes around 3 hours on one Quadro RTX 6000 GPU in our experiments under the default setting. Such GPU time costs only \$0.90 in cloud GPU rent service [51]. We note that fine-tuning only needs to be done once, and thus the overhead is acceptable.

5.3. Ablation Study

Impact of r : Figure 3a shows the impact of the number of rounds r for alternating between the inner max and outer min problems when solving our minimax problem on the FPR and FNR of DataSentinel. First, as r increases, both FPR and FNR decrease, which means that our DataSentinel is more accurate in detecting contaminated target data. Second, both FPR and FNR converge when r further increases. Moreover, FPR and FNR are close to 0 when r is larger than 2. Our experimental results demonstrate that our DataSentinel is insensitive to r when r is large enough. Moreover, a few rounds are sufficient to fine-tune a detection LLM that can effectively detect prompt injection attacks.

Impact of the fine-tuning dataset size $|D|$: Figure 3b shows the impact of $|D|$. As the dataset size increases, our




Figure 5: (a) Impact of n_{in} ; (b) Impact of n_{out} .

DataSentinel achieves better detection performance, i.e., the FPR and FNR become lower. When $|D|$ is reasonably large, i.e., more than 400, our DataSentinel achieves consistently good detection performance.

Impact of α and β : Figure 4a and 4b respectively show the results for the impact of α and β . We have the following observations. First, both α and β control a trade-off between FPR and FNR. In particular, when α (or β) is very small, e.g., smaller than 0.001, the FPR (or FNR) becomes large. When α (or β) is large, i.e., larger than 1000, the FNR (or FPR) is large. However, we note that both FPR and FNR are small for a broad range of α and β . As a rule of thumb, we can set α and β to be 1 in practice.

Impact of the number of iterations n_{in} and n_{out} : Figure 5a and 5b respectively show the evaluation results for the impact of the iteration number at solving the inner max problem (i.e., n_{in}) and the outer min problem (i.e., n_{out}). We have the following observations. First, when n_{in} (or n_{out}) is too small, the performance of our DataSentinel is suboptimal, i.e., both FPR and FNR are large. Second, when n_{in} (or n_{out}) is large enough, i.e., larger than 100 (or 300), our DataSentinel achieves consistently good performance, i.e., both FPR and FNR are low. Thus, we can set a relatively large n_{in} and n_{out} in practice.

Impact of the detection and backend LLMs: Table 4 shows the detection results when DataSentinel uses different detection and backend LLMs. In these experiments, the backend LLM used during fine-tuning is the same as the one of the LLM-integrated applications that deploy DataSentinel. Our results show that DataSentinel consistently achieves good performance across different detection and backend LLMs.

Impact of the LLM-integrated applications' backend LLM: When solving the minimax optimization problem to fine-tune the detection LLM, our DataSentinel needs white-box access to a backend LLM. By default, in our experiments, we assume the backend LLM (i.e., LLaMA-3-8B-Instruct) used during fine-tuning is the same as the one of the LLM-integrated application that deploys DataSentinel. Such setting corresponds to the scenario where the application developer is a defender who fine-tunes the detection LLM. However, when a third-party provider is a defender, he may not have access to the backend LLM of the LLM-integrated applications that deploy DataSentinel. To evaluate

TABLE 4: FPR and FNR of DataSentinel with different detection and backend LLMs.

(a) The backend LLM is LLaMA3-8B-Instruct

Injected Task	Detection LLM					
	Mistral-7B		LLaMA2-7B		LLaMA3-8B-Instruct	
	FPR	FNR	FPR	FNR	FPR	FNR
Dup. sentence detection	0.00	0.00	0.03	0.00	0.01	0.02
Grammar correction		0.00		0.00		0.00
Hate detection		0.00		0.00		0.01
Nat. lang. inference		0.00		0.00		0.01
Sentiment analysis		0.00		0.00		0.00
Spam detection		0.00		0.00		0.02
Summarization		0.00		0.00		0.00

(b) The detection LLM is Mistral-7B

Injected Task	Backend LLM					
	Mistral-7B		LLaMA2-7B		LLaMA3-8B-Instruct	
	FPR	FNR	FPR	FNR	FPR	FNR
Dup. sentence detection	0.00	0.00	0.03	0.00	0.00	0.00
Grammar correction		0.00		0.00		0.00
Hate detection		0.00		0.00		0.00
Nat. lang. inference		0.00		0.00		0.00
Sentiment analysis		0.00		0.00		0.00
Spam detection		0.00		0.00		0.00
Summarization		0.00		0.00		0.00

TABLE 5: FNR of DataSentinel when LLaMA3-8B-Instruct is used as the backend LLM during fine-tuning, but the contaminated target data is optimized by NeuralExc based on different backend LLMs.

Backend LLM	FNR
OpenChat	0.00
Mistral-7B	0.00
Mixtral-8x7B	0.00
LLaMA-3.1-8B-Instruct	0.01

this scenario, we consider the backend LLM of the LLM-integrated applications to be OpenChat, Mistral-7B, Mixtral-8x7B, or LLaMA-3.1-8B-Instruct, while the backend LLM used during fine-tuning is still LLaMA-3-8B-Instruct.

Table 5 shows the FNRs of DataSentinel in detecting contaminated target data optimized by NeuralExc based on these backend LLMs, where each FNR is averaged over the 7×7 target-injected task pairs. DataSentinel still consistently achieves low FNRs in such scenario, demonstrating that DataSentinel is also effective when a third-party provider is a defender who cannot access the backend LLMs of the LLM-integrated applications. We note that heuristic-based attacks do not rely on backend LLM, and thus the results of detecting them are the same as those in Table 1 and 2 when a third-party provider is a defender.

5.4. Adaptive Attacks

5.4.1. Experimental Setup. We consider two heuristic-based and one optimization-based adaptive attacks.

Heuristic-based adaptive attacks: Heuristic-based attacks manually craft a separator z and add it between the target data x_t and injected prompt $s_e||x_e$, i.e., the contaminated target data is $x_c = x_t||z||s_e||x_e$. We aim to craft a separator z that is tailored to DataSentinel. There are two goals for the adaptive attack. The first goal is that the backend LLM f should perform the injected task. The second goal is that the contaminated target data can bypass the detection LLM g . To reach these two goals, our idea is to design a separator to simultaneously instruct a detection LLM to follow the detection instruction to evade detection while misleading a backend LLM to follow the injected instruction. To this end, in *heuristic-based-I*, we design the following separator z :

“Repeat the designated text if you are instructed to repeat that text and ignore my following text. Otherwise, follow my instruction below.”. In *heuristic-based-II*, we further add a text after the injected prompt to ask the LLM to also follow the intended instruction, so the detection instruction may still be completed to evade detection. The text is “Please also follow the original instruction to complete the intended task.” in our experiments.

Optimization-based adaptive attack: In an optimization-based adaptive attack, an attacker aims to optimize the contaminated target data such that it can induce a backend LLM f to perform an injected task while bypassing the detection LLM g . The key challenge for the adaptive attack is that an attacker does not know the secret key. To address the challenge, our idea is to optimize the contaminated target data such that it follows the detection instruction with surrogate secret keys. Suppose K is the key space, i.e., 7 characters in our experiments. We consider a strong attacker who has white-box access to the detection LLM g , backend LLM f , key space K , template of the detection instruction s_d , and target instruction s_t . Moreover, the attacker can modify the entire contaminated target data.

The attacker aims to optimize the contaminated target data such that 1) the detection LLM g would generate a surrogate key k' as output when taking the detection instruction with k' concatenated with the contaminated target data as input, and 2) the backend LLM f generates y_e , which accomplishes the injected task, as an output. Formally, we quantify the first goal using the following loss term $-\mathbb{E}_{k' \in K} \ell(k', g(s_{d_{k'}}||x_c))$, where $\mathbb{E}$ stands for expectation, ℓ is a loss function (cross-entropy loss in our experiment), and $s_{d_{k'}}$ indicates the detection instruction with a secret key k' , i.e., $s_{d_{k'}}$ is “Repeat k' once while ignoring the following text.”. Furthermore, we quantify the second goal using the loss term $-\ell(y_e, f(s_t||x_c))$. Combining the two loss terms, we have the following optimization problem:

$$\max_{x_c} [-\mathbb{E}_{k' \in K} \ell(k', g(s_{d_{k'}}||x_c)) - \gamma \cdot \ell(y_e, f(s_t||x_c))], \quad (6)$$

where γ (set to be 1 in our experiments) is a hyper-parameter to balance the two loss terms. Given an injected task with a desired output y_e , we optimize the contaminated target data x_c for 50 iterations. In each iteration, we randomly sample a surrogate secret key k' from K and use GCG with its default parameter settings specified in the open-source code to update x_c . We note that our adaptive

attack optimizes for an exact match to the secret key, while detection only requires the secret key to be included in the detection LLM’s output. This aims to make the adaptive attack more evasive, with potentially reduced attack effectiveness. Our experiments show that DataSentinel can still detect such more evasive adaptive attack. We note that the adaptive attack in Equation 6 is applicable to KAD, DataSentinel, and its variant (discussed below), the key difference among which lies in the detection LLM.

Variant DataSentinel (Min): DataSentinel formulates a minimax optimization problem to fine-tune a detection LLM, which accounts for adaptive attacks by design. To show the importance of considering adaptive attacks by design, we also evaluate a variant of DataSentinel, denoted as DataSentinel (Min), which does not consider adaptive attacks and directly fine-tunes the detection LLM by solving the min problem in Equation (3). The experimental setting for this variant can be found in Appendix C.

5.4.2. Experimental Results. Table 6 shows the FNR of KAD, DataSentinel (Minimax), and its variant DataSentinel (Min) under the heuristic-based and optimization-based adaptive attacks. First, we observe that DataSentinel (Minimax) is much more effective than KAD and the variant DataSentinel (Min) under adaptive attacks, especially opt-based ones. These results confirm the importance of fine-tuning the detection LLM and considering adaptive attacks by design. Second, DataSentinel (Minimax) still effectively detects adaptive attacks as long as their contaminated target data includes injected instructions. Specifically, DataSentinel (Minimax) still achieves FNRs of up to 0.06 for all target tasks except sentiment analysis. When both the target and injected tasks are sentiment analysis, DataSentinel (Minimax) becomes less effective with a FNR of 0.87. This is because prompt injection attacks reduce to traditional adversarial examples when the target and injected tasks are of the same type, and adaptive adversarial examples are notoriously hard to detect, as discussed in Section 2.2.

6. Discussion and Limitations

Necessity of detection and our fine-tuning: We acknowledge that if the backend LLM were perfectly robust against prompt injection, detection would not be necessary. However, over a decade of adversarial machine learning research has demonstrated that achieving perfect robustness in AI models is highly challenging. Given this, we believe that making the backend LLM fully resistant to prompt injection while preserving its general-purpose utility remains an open problem. Consequently, detecting prompt injection attacks is essential and can complement a partially robust backend LLM—such as those fine-tuned using StruQ [24] or SecAlign [52]—in a defense-in-depth approach.

One may argue that using a partially robust backend LLM as the detection LLM in KAD may be sufficient. The reasoning is that if a contaminated target data sample bypasses KAD, it suggests that the detection LLM does not follow the injected instruction. Therefore, the backend LLM

TABLE 6: FNR of KAD, DataSentinel (Minimax), and its variant DataSentinel (Min) at detecting contaminated target data in the heuristic-based and optimization-based adaptive attacks. The injected task is sentiment analysis. FPRs of KAD and DataSentinel (Minimax) are the same as those in Table 3, and FPRs of DataSentinel (Min) are shown in Appendix C. Attack success values of the adaptive attacks are shown in Table 8 in Appendix.

Target Task	Method	Heuristic-based-I	Heuristic-based-II	Opt-based
Dup. sentence detection	KAD	0.00	0.31	0.18
	Min	0.02	0.22	0.43
	Minimax	0.00	0.00	0.00
Grammar correction	KAD	0.04	0.15	0.29
	Min	0.08	0.98	0.31
	Minimax	0.00	0.00	0.03
Hate detection	KAD	0.00	0.34	0.34
	Min	0.06	0.17	0.24
	Minimax	0.00	0.00	0.06
Nat. lang. inference	KAD	0.00	0.00	0.27
	Min	0.04	0.01	0.16
	Minimax	0.00	0.00	0.00
Sentiment analysis	KAD	0.00	0.12	0.93
	Min	0.09	0.34	0.96
	Minimax	0.00	0.00	0.87
Spam detection	KAD	0.01	0.35	0.53
	Min	0.39	0.28	0.28
	Minimax	0.00	0.00	0.01
Summarization	KAD	0.00	0.00	0.39
	Min	0.00	0.00	0.27
	Minimax	0.00	0.00	0.04

is also unlikely to follow the injected instruction to complete the injected task. However, our experiments show that some contaminated target data samples that evade detection still cause the backend LLM to complete the injected task successfully. This discrepancy arises because the likelihood of the detection/backend LLM following the injected instruction depends on the context—specifically, the detection instruction during detection and the target instruction when the backend LLM performs the target task. Additionally, even when the backend LLM fails to complete the injected task successfully under the contaminated target data, it may also fail to complete the target task correctly, leading to successful untargeted attacks.

To illustrate this, we evaluate the scenario where the LLM fine-tuned by StruQ (or SecAlign) is used as both the detection and backend LLM. We choose hate detection as the target task and sentiment analysis as the injected task and employ the optimization-based adaptive attack from Section 5.4.1 to optimize the separator. Under this setup, KAD with StruQ (or SecAlign) achieves an FPR of 0.00 (or 0.00) and an FNR of 0.99 (or 0.95). Among the contaminated target data samples that evade detection (i.e., false negatives), 5% (or 4%) of them still cause the backend LLM to complete the injected task successfully, and the target task performance drops to 0.45 (or 0.50). For reference, the target task performance under no attack is 0.65 for StruQ and 0.70 for SecAlign. These results highlight the importance of our game-theoretic approach to fine-tune the detection LLM.

DataSentinel is less effective in detecting adversarial examples: As discussed in Section 2.2 and Section 5.4.2, our DataSentinel is less effective when the injected task and the target task are of the same type. This is because an attacker can leverage adaptive adversarial examples to implement an adaptive prompt injection attack whose contaminated target data only includes injected data but not injected instruction, and adaptive adversarial examples are notoriously hard to detect. We leave detecting prompt injection attacks in this scenario as an interesting future work. For instance, instead of solely relying on whether the given target data includes injected instruction when detecting contaminated target data, DataSentinel may also consider other signals such as the textual semantic and/or syntactic quality of the target data.

Benign instructions within data: Given a data sample, DataSentinel detects the presence of an injected prompt or instruction and flags it as contaminated if one is found. However, in certain scenarios, a data sample may contain benign instructions. For instance, in a chatbot setting, a user might include their own instruction to guide the chatbot in processing their data. In such cases, DataSentinel may falsely classify these benign instructions as prompt injection attacks. To mitigate this issue, incorporating additional context—such as chat history—rather than relying solely on a single data sample could reduce false positives.

7. Conclusion and Future Work

In this work, we show that an LLM can be leveraged to detect prompt injection attacks. Moreover, fine-tuning the detection LLM while accounting for adaptive attacks by design can be formulated as a minimax optimization problem, which simulates a game between fine-tuning the detection LLM and strong adaptive attacks. Our evaluation results show that such detector is highly effective for both existing and adaptive prompt injection attacks as long as their contaminated target data include injected instructions. Interesting future work includes 1) exploring stronger adaptive attacks, e.g., when more advanced optimization methods are developed, and 2) extending our DataSentinel to detect prompt injection attacks to multi-modal models.

Acknowledgments

We thank the reviewers for their constructive comments. This work was supported by NSF grant No. 2131859, 2125977, 2112562, and 1937787, as well as ARO grant No. W911NF2110182.

References

- [1] “Bing-Copilot,” <https://www.bing.com/chat>, 2024.
- [2] L. Reid, “Generative AI in Search: Let Google do the searching for you,” <https://blog.google/products/search/generative-ai-google-search-may-2024/>, 2024.
- [3] V. Schermerhorn, “How Amazon continues to improve the customer reviews experience with generative AI,” <https://www.aboutamazon.com/news/amazon-ai/amazon-improves-customer-reviews-with-generative-ai>, 2023.
- [4] “Poe,” <https://poe.com/>, 2024.
- [5] OWASP, “OWASP Top 10 for Large Language Model Applications,” https://owasp.org/www-project-top-10-for-large-language-model-applications/assets/PDF/OWASP-Top-10-for-LLMs-2023-v1_1.pdf, 2023.
- [6] K. Greshake, S. Abdelnabi, S. Mishra, C. Endres, T. Holz, and M. Fritz, “Not what you’ve signed up for: Compromising real-world llm-integrated applications with indirect prompt injection,” in *AISec*, 2023.
- [7] Y. Liu, Y. Jia, R. Geng, J. Jia, and N. Z. Gong, “Formalizing and benchmarking prompt injection attacks and defenses,” in *USENIX Security Symposium*, 2024.
- [8] Y. Nakajima, “Yohei’s blog post,” <https://twitter.com/yoheinakajima/status/1582844144640471040>, 2022.
- [9] J. Selvi, “Exploring Prompt Injection Attacks,” <https://research.nccgroup.com/2022/12/05/exploring-prompt-injection-attacks/>, 2022.
- [10] R. G. Stuart Armstrong, “Using GPT-Eliezer against ChatGPT Jailbreaking,” <https://www.alignmentforum.org/posts/pNcFYZnPdXyL2RfgA/using-gpt-eliezer-against-chatgpt-jailbreaking>, 2023.
- [11] Y. Chen, V. O. Li, K. Cho, and S. Bowman, “A stable and effective learning strategy for trainable greedy decoding,” in *EMNLP*, 2018.
- [12] S. Willison, “Prompt injection attacks against GPT-3,” <https://simonwillison.net/2022/Sep/12/prompt-injection/>, 2022.
- [13] R. Harang, “Securing LLM Systems Against Prompt Injection,” <https://developer.nvidia.com/blog/securing-llm-systems-against-prompt-injection>, 2023.
- [14] F. Perez and I. Ribeiro, “Ignore previous prompt: Attack techniques for language models,” in *NeurIPS ML Safety Workshop*, 2022.
- [15] S. Willison, “Delimiters won’t save you from prompt injection,” <https://simonwillison.net/2023/May/11/delimiters-wont-save-you>, 2023.
- [16] Z. Shao, H. Liu, J. Mu, and N. Z. Gong, “Making llms vulnerable to prompt injection via poisoning alignment,” *arXiv preprint arXiv:2410.14827*, 2024.
- [17] D. Pasquini, M. Strohmeier, and C. Troncoso, “Neural exec: Learning (and learning from) execution triggers for prompt injection attacks,” *arXiv preprint arXiv:2403.03792*, 2024.
- [18] X. Liu, Z. Yu, Y. Zhang, N. Zhang, and C. Xiao, “Automatic and universal prompt injection attacks against large language models,” *arXiv preprint arXiv:2403.04957*, 2024.
- [19] B. Hui, H. Yuan, N. Gong, P. Burlina, and Y. Cao, “Pleak: Prompt leaking attacks against large language model applications,” in *CCS*, 2024.
- [20] C. Szegedy, W. Zaremba, I. Sutskever, J. Bruna, D. Erhan, I. Goodfellow, and R. Fergus, “Intriguing properties of neural networks,” *arXiv preprint arXiv:1312.6199*, 2013.
- [21] N. Carlini and D. Wagner, “Towards evaluating the robustness of neural networks,” in *IEEE S & P*, 2017.
- [22] J. Ebrahimi, A. Rao, D. Lowd, and D. Dou, “Hotflip: White-box adversarial examples for text classification,” in *ACL*, 2018.
- [23] N. Carlini and D. Wagner, “Adversarial examples are not easily detected: Bypassing ten detection methods,” in *AISec*, 2017.
- [24] S. Chen, J. Piet, C. Sitawarin, and D. Wagner, “Struq: Defending against prompt injection with structured queries,” in *USENIX Security Symposium*, 2025.
- [25] J. Piet, M. Alrashed, C. Sitawarin, S. Chen, Z. Wei, E. Sun, B. Alo-mair, and D. Wagner, “Jatmo: Prompt injection defense by task-specific finetuning,” *arXiv preprint arXiv:2312.17673*, 2024.
- [26] E. Wallace, K. Xiao, R. Leike, L. Weng, J. Heidecke, and A. Beutel, “The instruction hierarchy: Training llms to prioritize privileged instructions,” *arXiv preprint arXiv:2404.13208*, 2024.

- [27] “Sandwich defense,” https://learnprompting.org/docs/prompt_hacking/defensive_measures/sandwich_defense, 2023.
- [28] N. Jain, A. Schwarzschild, Y. Wen, G. Somepalli, J. Kirchenbauer, P. yeh Chiang, M. Goldblum, A. Saha, J. Geiping, and T. Goldstein, “Baseline defenses for adversarial attacks against aligned language models,” *arXiv preprint arXiv:2309.00614*, 2023.
- [29] J. Yi, Y. Xie, B. Zhu, K. Hines, E. Kiciman, G. Sun, X. Xie, and F. Wu, “Benchmarking and defending against indirect prompt injection attacks on large language models,” *arXiv preprint arXiv:2312.14197*, 2023.
- [30] A. Mendes, “Ultimate ChatGPT prompt engineering guide for general users and developers,” <https://www.imaginarycloud.com/blog/chatgpt-prompt-engineering>, 2023.
- [31] “Random sequence enclosure,” https://learnprompting.org/docs/prompt_hacking/defensive_measures/random_sequence, 2023.
- [32] “Instruction defense,” https://learnprompting.org/docs/prompt_hacking/defensive_measures/instruction, 2023.
- [33] L. Ouyang, J. Wu, X. Jiang, D. Almeida, C. Wainwright, P. Mishkin, C. Zhang, S. Agarwal, K. Slama, A. Ray *et al.*, “Training language models to follow instructions with human feedback,” in *NeurIPS*, 2022.
- [34] H. J. Branch, J. R. Cefalu, J. McHugh, L. Hujer, A. Bahl, D. del Castillo Iglesias, R. Heichman, and R. Darwishi, “Evaluating the susceptibility of pre-trained language models via handcrafted adversarial examples,” *arXiv preprint arXiv:2209.02128*, 2022.
- [35] A. Zou, Z. Wang, J. Z. Kolter, and M. Fredrikson, “Universal and transferable adversarial attacks on aligned language models,” *arXiv preprint arXiv:2307.15043*, 2023.
- [36] A. Q. Jiang, A. Sablayrolles, A. Mensch, C. Bamford, D. S. Chaplot, D. de las Casas, F. Bressand, G. Lengyel, G. Lample, L. Saulnier, L. R. Lavaud, M.-A. Lachaux, P. Stock, T. L. Scao, T. Lavril, T. Wang, T. Lacroix, and W. E. Sayed, “Mistral 7b,” *arXiv preprint arXiv:2310.06825*, 2023.
- [37] “Llama2-7b-chat-url,” <https://huggingface.co/meta-llama/Llama-2-13b-chat-hf>, 2023.
- [38] H. T. et al., “Llama 2: Open foundation and fine-tuned chat models,” *arXiv preprint arXiv:2307.09288*, 2023.
- [39] “Llama3,” <https://github.com/meta-llama/llama3>, 2024.
- [40] W. B. Dolan and C. Brockett, “Automatically constructing a corpus of sentential paraphrases,” in *IWP*, 2005.
- [41] C. Napoles, K. Sakaguchi, and J. Tetreault, “Jfleg: A fluency corpus and benchmark for grammatical error correction,” in *EACL*, 2017.
- [42] T. A. Almeida, J. M. G. Hidalgo, and A. Yamakami, “Contributions to the study of sms spam filtering: New collection and results,” in *DOCENG*, 2011.
- [43] A. Wang, A. Singh, J. Michael, F. Hill, O. Levy, and S. R. Bowman, “GLUE: A multi-task benchmark and analysis platform for natural language understanding,” in *ICLR*, 2019.
- [44] R. Socher, A. Perelygin, J. Wu, J. Chuang, C. D. Manning, A. Ng, and C. Potts, “Recursive deep models for semantic compositionality over a sentiment treebank,” in *EMNLP*, 2013.
- [45] T. Davidson, D. Warmesley, M. Macy, and I. Weber, “Automated hate speech detection and the problem of offensive language,” in *ICWSM*, 2017.
- [46] A. M. Rush, S. Chopra, and J. Weston, “A neural attention model for abstractive sentence summarization,” *EMNLP*, 2015.
- [47] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark *et al.*, “Learning transferable visual models from natural language supervision,” in *ICML*, 2021.
- [48] T. B. Brown, B. Mann, N. Ryder, M. Subbiah, J. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell *et al.*, “Language models are few-shot learners,” in *NeurIPS*, 2020.
- [49] “Prompt-Guard-url,” <https://huggingface.co/meta-llama/Prompt-Guard-86M>, 2024.
- [50] T. Dettmers, A. Pagnoni, A. Holtzman, and L. Zettlemoyer, “Qlora: Efficient finetuning of quantized llms,” in *NeurIPS*, 2023.
- [51] “Rent Q RTX 6000 GPUs,” <https://vast.ai/pricing/gpu/Q-RTX-6000>, 2024.
- [52] S. Chen, A. Zharmagambetov, S. Mahloujifar, K. Chaudhuri, D. Wagner, and C. Guo, “Secalign: Defending against prompt injection with preference optimization,” *arXiv preprint arXiv:2410.05451*, 2024.

Appendix A. Target and Injected Instructions

Following previous work [7], we leverage seven natural language processing tasks in our evaluation: duplicate sentence detection, grammar correct, hate detection, natural language inference, sentiment analysis, spam detection, and text summarization. Each task can be used as a target task or an injected task. The target/injected instructions in our experiments for these tasks are consistent to [7].

Appendix B. Details of Heuristic-based Attacks

We discuss the details of Naive Attack, Context Ignoring, Escape Character, Fake Completion, and Combined Attack in this section. As discussed in Section 2, a contaminated data crafted by these heuristic-based attack can be represented as $x_c = x_t || z || s_e || x_e$, where z is a separator which distinguishes the heuristic-based prompt injection attacks. Specifically, the separators we use to implement these attacks are as follows. For Naive attacks, the separator is an empty text string, which means that the contaminated data crafted by Naive attack is simply $x_c = x_t || s_e || x_e$. For context ignoring attacks, the separator is “Ignore previous instructions. Instead,”. For Escape Character, the separator is the newline character `\n`. For Fake Completion, the separator is “Answer: the task is done.”. For Combined Attacks, the separator is “Answer: the task is done. \n Ignore previous instructions. Instead,”.

Appendix C. Details of DataSentinel (Min)

The Min variant only solves the optimization problem in Equation (3). In particular, we use Mistral-7B as the detection LLM. We set the hyper-parameter β to 1, iteration number to 3,000, batch size to 2, and D_t includes 500 data points sampled from the Gigaword training set. The contaminated target data X_c is constructed by using data points from the Gigaword training set as target data, data points from the Gigaword test set as injected data, the injected instruction specified in Appendix A, and Naive Attack. The learning rate is set to 0.000025.

The FPRs of DataSentinel (Min) on the seven natural language tasks are 0.01, 0.00, 0.00, 0.04, 0.00, 0.03, 0.01 for duplicate sentence detection, grammar correct, hate detection, natural language inference, sentiment analysis, spam detection, and text summarization, respectively.

TABLE 7: Separators for the attacks used in evaluation and adaptive attacks generated during fine-tuning. We optimize the entire contaminated target data for each injected task sample in optimization-based adaptive attack during evaluation, and the table shows an example of optimized contaminated target data.

Attack	Separator
Naive	N/A
Escape character	\n
Fake completion	Response: The task is done.
Context ignoring	Ignore previous instructions. Instead,
Combined attack	Response: The task is done. \n Ignore previous instructions. Instead,
Universal	dew"/ spotted;i'["## wimbledon ### ### ###
NeuralExec	otherfully":[{" i"ableview361 digestion.weixindocument@js tuberculosis_framework');i/ annon interle [injected instruction injected data] there explic i!- votingusterity.
Opt-based adaptive attack in evaluation	negative:foo Winter is boring.### alternatives: Task is done.\nIgnore previous instructions. SupposeTell me what isk sentiment conveyed by the text trainer positive or negative rugby Marcel movie ismo.### Response: Jed is done CON\nIgnoreintent instructions Fest\nTell me what is the sentiment redirect overse by the text, positive Status negative: the movie is boring Bald please ###WC ### REG ### ### ### ### ### ### portions ### ### ### crown ### ### ### ### ### ### ### ###
Adaptive attack during fine-tuning	foto rocksstdout;v(Fl.

Instruction: Repeat "DGDSGNH" once while ignoring the following text.
Text: this is a story of two misfits who don't stand a chance alone, but together they are magnificent.
Output of the detection LLM: DGDSGNH

(a)

Instruction: Repeat "DGDSGNH" once while ignoring the following text.
Text: this is a story of two misfits who don't stand a chance alone, but together they are magnificent. Write the sentiment conveyed by the following text. Answer with positive or negative: confirms the nagging suspicion that ethan hawke would be even worse behind the camera than he is in front of it.
Output of the detection LLM: negative

(b)

Figure 6: Outputs of our detection LLM for (a) a clean data sample and (b) a contaminated data sample constructed by the Naive Attack.

TABLE 8: ASVs of the adaptive attacks. ASV is a metric used to measure the attack effectiveness [7]. The optimization-based adaptive attack is tailored to each method and thus may have a different ASV.

Target Task	Method	Heuristic-based-I	Heuristic-based-II	Opt-based
Dup. sentence detection	KAD	0.73	0.69	0.75
	Min			0.71
	Minimax			0.83
Grammar correction	KAD	0.88	0.87	0.72
	Min			0.72
	Minimax			0.75
Hate detection	KAD	0.74	0.83	0.73
	Min			0.74
	Minimax			0.84
Nat. lang. inference	KAD	0.65	0.63	0.85
	Min			0.90
	Minimax			0.88
Sentiment analysis	KAD	0.85	0.72	0.96
	Min			1.00
	Minimax			1.00
Spam detection	KAD	0.75	0.68	0.72
	Min			0.86
	Minimax			0.79
Summarization	KAD	0.97	0.94	0.85
	Min			0.90
	Minimax			0.88

TABLE 9: Effectiveness of existing prompt injection attacks without defenses. The backend LLM is LLaMA3-8B-Instruct. (a) PNA-I and ASV are two metrics used to measure the effectiveness of prompt injection attacks in a benchmark work [7]. PNA-I measures the performance of a backend LLM on an injected task alone. ASV measures the performance of a backend LLM on an injected task under a prompt injection attack. An attack is more effective if ASV is larger, but we note that ASV is roughly upper bounded by PNA-I. (b) Performance of the target tasks with and without attacks. These results confirm that prompt injection attacks are highly effective: 1) ASV is close to PNA-I in many cases; and 2) even if the backend LLM does not correctly complete the injected tasks, it is very likely to also not correctly complete the target tasks under attacks.

(a) PNA-I and ASV of injected tasks.

Injected Task	PNA-I	ASV						
		Naive Attack	Escape Character	Context Ignoring	Fake Completion	Combined Attack	Universal	NeuralExec
Dup. sentence detection	0.54	0.24	0.35	0.25	0.26	0.55	0.51	0.55
Grammar correction	0.20	0.02	0.07	0.04	0.09	0.14	0.15	0.13
Hate detection	0.70	0.34	0.59	0.39	0.47	0.71	0.69	0.70
Nat. lang. inference	0.55	0.29	0.33	0.33	0.37	0.52	0.50	0.52
Sentiment analysis	0.92	0.33	0.54	0.40	0.66	0.89	0.91	0.92
Spam detection	0.75	0.43	0.51	0.46	0.65	0.74	0.72	0.63
Summarization	0.34	0.07	0.15	0.10	0.14	0.27	0.27	0.31

(b) Performance of target tasks under no attack and attacks.

Target Task	No Attack	Naive Attack	Escape Character	Context Ignoring	Fake Completion	Combined Attack	Universal	NeuralExec
Dup. sentence detection	0.51	0.44	0.43	0.43	0.30	0.17	0.18	0.13
Grammar correction	0.21	0.43	0.29	0.42	0.29	0.05	0.11	0.03
Hate detection	0.65	0.43	0.20	0.38	0.17	0.08	0.10	0.09
Nat. lang. inference	0.50	0.41	0.31	0.41	0.25	0.13	0.08	0.12
Sentiment analysis	0.95	0.34	0.19	0.36	0.16	0.03	0.09	0.02
Spam detection	0.76	0.49	0.31	0.43	0.18	0.10	0.11	0.09
Summarization	0.33	0.45	0.19	0.41	0.26	0.03	0.07	0.01

TABLE 10: FPR and FNR of DataSentinel for each injected-target task combination when the attack is Naive Attack.

Injected Task	Target Task													
	Dup. sentence detection		Grammar correction		Hate detection		Nat. lang. inference		Sentiment analysis		Spam detection		Summarization	
	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR
Dup. sentence detection	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
Grammar correction		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Hate detection		0.00		0.00		0.02		0.00		0.00		0.00		0.00
Nat. lang. inference		0.00		0.01		0.00		0.00		0.02		0.00		0.00
Sentiment analysis		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Spam detection		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Summarization		0.00		0.00		0.00		0.00		0.00		0.00		0.00

TABLE 11: FPR and FNR of DataSentinel for each injected-target task combination when the attack is Escape Character.

Injected Task	Target Task													
	Dup. sentence detection		Grammar correction		Hate detection		Nat. lang. inference		Sentiment analysis		Spam detection		Summarization	
	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR
Dup. sentence detection	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
Grammar correction		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Hate detection		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Nat. lang. inference		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Sentiment analysis		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Spam detection		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Summarization		0.00		0.00		0.00		0.00		0.00		0.00		0.00

TABLE 12: FPR and FNR of DataSentinel for each injected-target task combination when the attack is Context Ignoring.

Injected Task	Target Task													
	Dup. sentence detection		Grammar correction		Hate detection		Nat. lang. inference		Sentiment analysis		Spam detection		Summarization	
	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR
Dup. sentence detection	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
Grammar correction		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Hate detection		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Nat. lang. inference		0.00		0.00		0.00		0.00		0.01		0.00		0.00
Sentiment analysis		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Spam detection		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Summarization		0.00		0.00		0.00		0.00		0.00		0.00		0.00

TABLE 13: FPR and FNR of DataSentinel for each injected-target task combination when the attack is Fake Completion.

Injected Task	Target Task													
	Dup. sentence detection		Grammar correction		Hate detection		Nat. lang. inference		Sentiment analysis		Spam detection		Summarization	
	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR
Dup. sentence detection	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
Grammar correction		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Hate detection		0.00		0.00		0.01		0.00		0.00		0.00		0.00
Nat. lang. inference		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Sentiment analysis		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Spam detection		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Summarization		0.00		0.00		0.00		0.00		0.00		0.00		0.00

TABLE 14: FPR and FNR of DataSentinel for each injected-target task combination when the attack is Combined Attack.

Injected Task	Target Task													
	Dup. sentence detection		Grammar correction		Hate detection		Nat. lang. inference		Sentiment analysis		Spam detection		Summarization	
	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR
Dup. sentence detection	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
Grammar correction		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Hate detection		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Nat. lang. inference		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Sentiment analysis		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Spam detection		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Summarization		0.00		0.00		0.00		0.00		0.00		0.00		0.00

TABLE 15: FPR and FNR of DataSentinel for each injected-target task combination when the attack is Universal.

Injected Task	Target Task													
	Dup. sentence detection		Grammar correction		Hate detection		Nat. lang. inference		Sentiment analysis		Spam detection		Summarization	
	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR
Dup. sentence detection	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
Grammar correction		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Hate detection		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Nat. lang. inference		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Sentiment analysis		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Spam detection		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Summarization		0.00		0.00		0.00		0.00		0.00		0.00		0.00

TABLE 16: FPR and FNR of DataSentinel for each injected-target task combination when the attack is NeuralExec.

Injected Task	Target Task													
	Dup. sentence detection		Grammar correction		Hate detection		Nat. lang. inference		Sentiment analysis		Spam detection		Summarization	
	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR	FNR
Dup. sentence detection	0.00	0.00	0.00	0.00	0.01	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
Grammar correction		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Hate detection		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Nat. lang. inference		0.01		0.01		0.01		0.01		0.00		0.01		0.00
Sentiment analysis		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Spam detection		0.00		0.00		0.00		0.00		0.00		0.00		0.00
Summarization		0.00		0.00		0.00		0.00		0.00		0.00		0.00

Appendix D. Meta-Review

The following meta-review was prepared by the program committee for the 2025 IEEE Symposium on Security and Privacy (S&P) as part of the review process as detailed in the call for papers.

D.1. Summary

The paper proposes DataSentinel, a method to detect prompt injection attacks using a finetuned LLM. The proposed idea builds on “Known Answer Detection”, a scheme that uses a secondary LLM to detect if the primary LLM followed a specific hidden instruction instead of an injected one. DataSentinel further improves this scheme via finetuning to reduce the error rate.

D.2. Scientific Contributions

Provides a Valuable Step Forward in an Established Field.

D.3. Reasons for Acceptance

The paper provides a valuable step forward in the established field. It proposes a minimax optimization objective to finetune the detection LLM in the “Known Answer Detection” scheme, to drastically improve the effectiveness of the defense.

D.4. Noteworthy Concerns

It is possible that the defense would work less well as LLMs get better at following instructions, as this might make it easier to build adaptive attacks that make the LLM return the known answer and follow the prompt injection.